

Introduction

In this notebook, we conduct an end-to-end analysis of the **BankChurners** dataset (15 numerical & 8 categorical features) to understand and predict customer attrition in a competitive banking environment. Our objectives are to:

1. **Define the problem** — quantify the impact of high churn (revenue loss, weakened brand loyalty, costly replacements).
2. **Prepare the data** — remove irrelevant columns, handle outliers via IQR, and convert `Attrition_Flag` into a binary `Churn` label.
3. **Explore patterns (EDA)** — visualise distributions and joint relationships across demographics, transaction volume, utilisation ratios, and tenure.
4. **Rank features** — use correlation, F-score (ANOVA), and mutual information to identify top predictors such as `Total_Trans_Ct`, `Total_Trans_Amt`, and `Avg_Utilization_Ratio`.
5. **Build models** — train and compare Logistic Regression, Decision Tree, Random Forest, AdaBoost, XGBoost, SVM, and MLP to find the best churn-prediction engine.
6. **Derive insights & actions** — translate model outputs into targeted retention strategies (workshops, credit-rebuilding, personalised perks).

1. Setup

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import chi2_contingency
import matplotlib.patches as mpatches
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier
from sklearn.metrics import (classification_report, confusion_matrix,
                             ConfusionMatrixDisplay, roc_curve,
                             roc_auc_score,
                             precision_recall_fscore_support)

plt.rcParams['figure.figsize'] = (10, 6)

import warnings
warnings.filterwarnings('ignore')
```

```
print("Setup complete.")
```

Setup complete.

2. Data Loading

Ensure the **BankChurners.csv** file is in the same directory as this notebook.

```
# Load the dataset
df = pd.read_csv("BankChurners.csv")
```

```
# Display the first 5 rows
df.head()
```

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count
0	768805383	Existing Customer	45	M	3
1	818770008	Existing Customer	49	F	5
2	713982108	Existing Customer	51	M	3
3	769911858	Existing Customer	40	F	4
4	709106358	Existing Customer	40	M	3

	Education_Level	Marital_Status	Income_Category	Card_Category
0	High School	Married	\$60K - \$80K	Blue
1	Graduate	Single	Less than \$40K	Blue
2	Graduate	Married	\$80K - \$120K	Blue
3	High School	Unknown	Less than \$40K	Blue
4	Uneducated	Married	\$60K - \$80K	Blue

	Months_on_book	...	Credit_Limit	Total_Revolving_Bal
0	39	...	12691.0	777
1	44	...	8256.0	864
2	36	...	3418.0	0
3	34	...	3313.0	2517
4	21	...	4716.0	0

	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct
0	11914.0
1	7392.0
2	3418.0
3	796.0
4	4716.0

0	1.335	1144	42
1.625			
1	1.541	1291	33
3.714			
2	2.594	1887	20
2.333			
3	1.405	1171	20
2.333			
4	2.175	816	28
2.500			

	Avg_Utilization_Ratio \
0	0.061
1	0.105
2	0.000
3	0.760
4	0.000

Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_1 \

0	0.000093
1	0.000057
2	0.000021
3	0.000134
4	0.000022

Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_2

0	0.99991
1	0.99994
2	0.99998
3	0.99987
4	0.99998

[5 rows x 23 columns]

3. Data Cleaning

3.1 Drop Unwanted Columns

We remove the following columns:

- CLIENTNUM
- Any columns generated by a Naive Bayes model (i.e., those starting with Naive_Bayes_Classifier_...)

```
# Define columns to drop
columns_to_drop = [
    'CLIENTNUM',

    'Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_1',

    'Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_2'
]

# Drop the specified columns
df.drop(columns=columns_to_drop, inplace=True, errors='ignore')

print("Columns removed.")
df.head()
```

Columns removed.

	Attrition_Flag	Customer_Age	Gender	Dependent_count	
0	Existing Customer	45	M	3	High School
1	Existing Customer	49	F	5	Graduate
2	Existing Customer	51	M	3	Graduate
3	Existing Customer	40	F	4	High School
4	Existing Customer	40	M	3	Uneducated

	Marital_Status	Income_Category	Card_Category	Months_on_book	
0	Married	\$60K - \$80K	Blue	39	
1	Single	Less than \$40K	Blue	44	
2	Married	\$80K - \$120K	Blue	36	
3	Unknown	Less than \$40K	Blue	34	
4	Married	\$60K - \$80K	Blue	21	

Total_Relationship_Count	Months_Inactive_12_mon
--------------------------	------------------------

Contacts_Count_12_mon	\		
0	5		1
3			
1	6		1
2			
2	4		1
0			
3	3		4
1			
4	5		1
0			

Credit_Limit	Total_Revolving_Bal	Avg_Open_To_Buy
Total_Amt_Chng_Q4_Q1	\	
0	12691.0	777
1.335		11914.0
1	8256.0	864
1.541		7392.0
2	3418.0	0
2.594		3418.0
3	3313.0	2517
1.405		796.0
4	4716.0	0
2.175		4716.0

Total_Trans_Amt	Total_Trans_Ct	Total_Ct_Chng_Q4_Q1
Avg_Utilization_Ratio		
0	1144	42
0.061		1.625
1	1291	33
0.105		3.714
2	1887	20
0.000		2.333
3	1171	20
0.760		2.333
4	816	28
0.000		2.500

3.2 Checking for Missing Values

We check for any missing values that might need to be addressed.

```
# Check for missing values
missing_values = df.isnull().sum()
print("Missing Values in each column:")
print(missing_values)
```

```
Missing Values in each column:
Attrition_Flag          0
```

```

Customer_Age      0
Gender            0
Dependent_count   0
Education_Level   0
Marital_Status    0
Income_Category   0
Card_Category     0
Months_on_book    0
Total_Relationship_Count  0
Months_Inactive_12_mon  0
Contacts_Count_12_mon  0
Credit_Limit     0
Total_Revolving_Bal  0
Avg_Open_To_Buy   0
Total_Amt_Chng_Q4_Q1  0
Total_Trans_Amt    0
Total_Trans_Ct     0
Total_Ct_Chng_Q4_Q1  0
Avg_Utilization_Ratio  0
dtype: int64

```

3.3 Create Numeric Churn Column

We keep the original `Attrition_Flag` for distribution plots. Then we create a numeric `Churn` column:

- Existing Customer → 0
- Attrited Customer → 1

```

# Map 'Attrition_Flag' to a binary numeric column 'Churn'
df['Churn'] = df['Attrition_Flag'].map({'Existing Customer': 0,
                                         'Attrited Customer': 1})

print("Unique values in Attrition_Flag:",
      df['Attrition_Flag'].unique())
print("Churn column created.")
df[['Attrition_Flag', 'Churn']].head()

```

```

Unique values in Attrition_Flag: ['Existing Customer' 'Attrited Customer']
Churn column created.

```

	Attrition_Flag	Churn
0	Existing Customer	0
1	Existing Customer	0
2	Existing Customer	0
3	Existing Customer	0
4	Existing Customer	0

4. Exploratory Data Analysis (EDA)

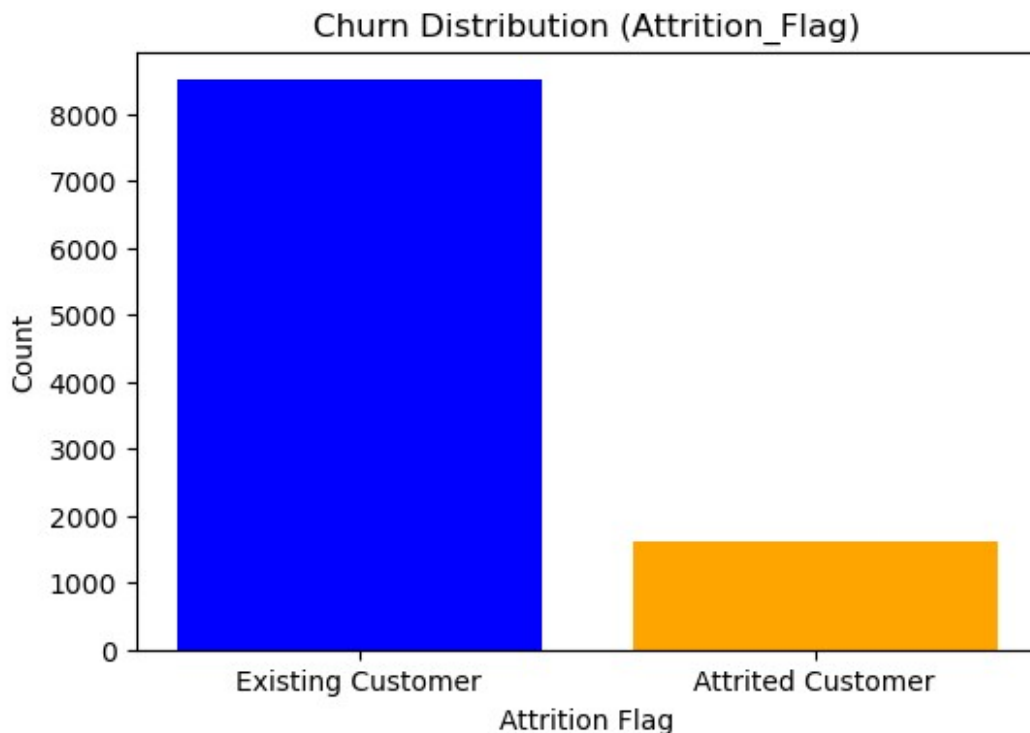
4.1 Distribution of Attrition_Flag

We examine how many customers are attrited vs. existing.

```
# Check distribution of Attrition_Flag
if 'Attrition_Flag' in df.columns:
    churn_counts = df['Attrition_Flag'].value_counts()
    print(churn_counts)

    # Bar plot
    plt.figure(figsize=(6,4))
    plt.bar(churn_counts.index, churn_counts.values,
            color=['blue', 'orange'])
    plt.title('Churn Distribution (Attrition_Flag)')
    plt.xlabel('Attrition Flag')
    plt.ylabel('Count')
    plt.show()
else:
    print("'Attrition_Flag' column not found.")

Attrition_Flag
Existing Customer    8500
Attrited Customer    1627
Name: count, dtype: int64
```



4.2 Outlier Removal

We remove outliers from all numeric columns **except** for Churn, using the IQR method.

```
# Create a copy for outlier removal
df_clean = df.copy()

# Identify numeric columns, excluding 'Churn'
numeric_cols =
df_clean.select_dtypes(include=[np.number]).columns.tolist()
if 'Churn' in numeric_cols:
    numeric_cols.remove('Churn')

# Remove outliers using IQR
for col in numeric_cols:
    Q1 = df_clean[col].quantile(0.25)
    Q3 = df_clean[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    df_clean = df_clean[(df_clean[col] >= lower_bound) &
(df_clean[col] <= upper_bound)]

print("Original shape:", df.shape)
print("Cleaned shape:", df_clean.shape)

# Update df to the cleaned version
df = df_clean.copy()

Original shape: (10127, 21)
Cleaned shape: (6463, 21)
```

4.3 Data Overview (Post-Cleaning)

```
# View basic info after cleaning
df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 6463 entries, 10 to 10125
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Attrition_Flag        6463 non-null   object
 1   Customer_Age          6463 non-null   int64
 2   Gender                6463 non-null   object
 3   Dependent_count       6463 non-null   int64
 4   Education_Level       6463 non-null   object
 5   Marital_Status        6463 non-null   object
 6   Income_Category       6463 non-null   object
 7   Card_Category         6463 non-null   object
```


8	Months_on_book	6463	non-null	int64
9	Total_Relationship_Count	6463	non-null	int64
10	Months_Inactive_12_mon	6463	non-null	int64
11	Contacts_Count_12_mon	6463	non-null	int64
12	Credit_Limit	6463	non-null	float64
13	Total_Revolving_Bal	6463	non-null	int64
14	Avg_Open_To_Buy	6463	non-null	float64
15	Total_Amt_Chng_Q4_Q1	6463	non-null	float64
16	Total_Trans_Amt	6463	non-null	int64
17	Total_Trans_Ct	6463	non-null	int64
18	Total_Ct_Chng_Q4_Q1	6463	non-null	float64
19	Avg_Utilization_Ratio	6463	non-null	float64
20	Churn	6463	non-null	int64

dtypes: float64(5), int64(10), object(6)

memory usage: 1.1+ MB

Summary statistics (including categorical columns)

df.describe(include='all')

	Attrition_Flag	Customer_Age	Gender	Dependent_count	\
count	6463	6463.000000	6463	6463.000000	
unique	2	NaN	2	NaN	
top	Existing Customer	NaN	F	NaN	
freq	5380	NaN	3988	NaN	
mean	NaN	46.384032	NaN	2.389138	
std	NaN	7.506139	NaN	1.272785	
min	NaN	26.000000	NaN	0.000000	
25%	NaN	41.000000	NaN	1.000000	
50%	NaN	46.000000	NaN	2.000000	
75%	NaN	52.000000	NaN	3.000000	
max	NaN	68.000000	NaN	5.000000	

	Education_Level	Marital_Status	Income_Category	Card_Category	\
count	6463	6463	6463	6463	
unique	7	4	6	4	
top	Graduate	Married	Less than \$40K	Blue	
freq	2006	2994	2677	6341	
mean	NaN	NaN	NaN	NaN	
std	NaN	NaN	NaN	NaN	
min	NaN	NaN	NaN	NaN	
25%	NaN	NaN	NaN	NaN	
50%	NaN	NaN	NaN	NaN	
75%	NaN	NaN	NaN	NaN	
max	NaN	NaN	NaN	NaN	

	Months_on_book	Total_Relationship_Count	...
Contacts_Count_12_mon	\		
count	6463.000000	6463.000000	...
6463.000000			
unique	NaN	NaN	...

NaN		
top	NaN	NaN ...
NaN		
freq	NaN	NaN ...
NaN		
mean	35.951106	3.939192 ...
2.497138		
std	7.036630	1.514048 ...
0.930392		
min	18.000000	1.000000 ...
1.000000		
25%	32.000000	3.000000 ...
2.000000		
50%	36.000000	4.000000 ...
3.000000		
75%	40.000000	5.000000 ...
3.000000		
max	53.000000	6.000000 ...
4.000000		

	Credit_Limit	Total_Revolving_Bal	Avg_Open_To_Buy \
count	6463.000000	6463.000000	6463.000000
unique	NaN	NaN	NaN
top	NaN	NaN	NaN
freq	NaN	NaN	NaN
mean	5084.608170	1140.233792	3944.374377
std	3921.874545	817.255076	3934.216696
min	1438.300000	0.000000	10.000000
25%	2264.000000	0.000000	1018.000000
50%	3359.000000	1250.000000	2229.000000
75%	6841.000000	1766.500000	5712.500000
max	19099.000000	2517.000000	16582.000000

	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct \
count	6463.000000	6463.000000	6463.000000
unique	NaN	NaN	NaN
top	NaN	NaN	NaN
freq	NaN	NaN	NaN
mean	0.728738	3593.651710	62.892155
std	0.162697	1568.378776	19.472826
min	0.293000	510.000000	10.000000
25%	0.620000	2216.000000	46.000000
50%	0.721000	3908.000000	67.000000
75%	0.836000	4581.000000	79.000000
max	1.193000	8454.000000	113.000000

	Total_Ct_Chng_Q4_Q1	Avg_Utilization_Ratio	Churn
count	6463.000000	6463.000000	6463.000000
unique	NaN	NaN	NaN
top	NaN	NaN	NaN

	NaN	NaN	NaN
freq			
mean	0.687565	0.320961	0.167569
std	0.181285	0.283666	0.373512
min	0.207000	0.000000	0.000000
25%	0.569000	0.000000	0.000000
50%	0.689000	0.260000	0.000000
75%	0.808000	0.570000	0.000000
max	1.182000	0.995000	1.000000

[11 rows x 21 columns]

4.4 Categorical Analysis

We analyze the average churn rate for key categorical variables.

```

categorical_cols = ["Gender", "Education_Level", "Marital_Status",
                    "Income_Category", "Card_Category"]

for col in categorical_cols:
    if col in df.columns:
        print(f"=== Analysis for {col} ===")
        churn_rate = df.groupby(col)['Churn'].mean()
        print("Churn Rate by Category:")
        print(churn_rate)

        # Optional: Chi-square test
        contingency_table = pd.crosstab(df[col], df['Churn'])
        chi2, p, dof, expected = chi2_contingency(contingency_table)
        print(f"\nChi-square statistic: {chi2:.4f}")
        print(f"p-value: {p:.4f}")
        print(f"Degrees of freedom: {dof}")
        print("-" * 50)


        # Plot bar chart
        plt.figure(figsize=(6, 4))
        churn_rate.plot(kind='bar', color='skyblue',
                        edgecolor='black')
        plt.title(f"Average Churn Rate by {col}")
        plt.ylabel("Churn Rate")
        plt.xticks(rotation=45)
        plt.tight_layout()
        plt.show()
    else:
        print(f"Column {col} not found in DataFrame.")

=== Analysis for Gender ===
Churn Rate by Category:
Gender
F      0.178536

```

```
M      0.149899
Name: Churn, dtype: float64
```

```
Chi-square statistic: 8.7742
p-value: 0.0031
Degrees of freedom: 1
```



```
=== Analysis for Education_Level ===
```

```
Churn Rate by Category:
```

```
Education_Level
```

```
College      0.158133
```

```
Doctorate    0.230769
```

```
Graduate     0.166002
```

```
High School  0.148148
```

```
Post-Graduate 0.195652
```

```
Uneducated   0.164871
```

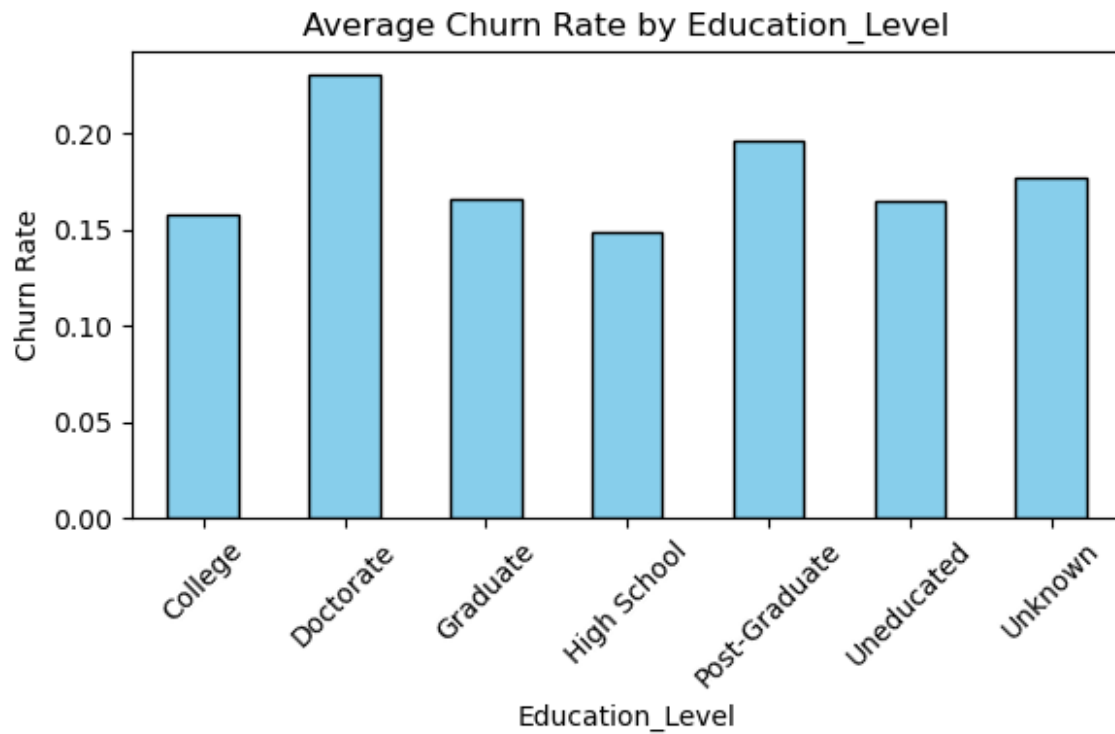
```
Unknown      0.177215
```

```
Name: Churn, dtype: float64
```

```
Chi-square statistic: 15.0267
```

```
p-value: 0.0201
```

```
Degrees of freedom: 6
```



=== Analysis for Marital_Status ===

Churn Rate by Category:

Marital_Status

Divorced 0.173362

Married 0.163995

Single 0.168526

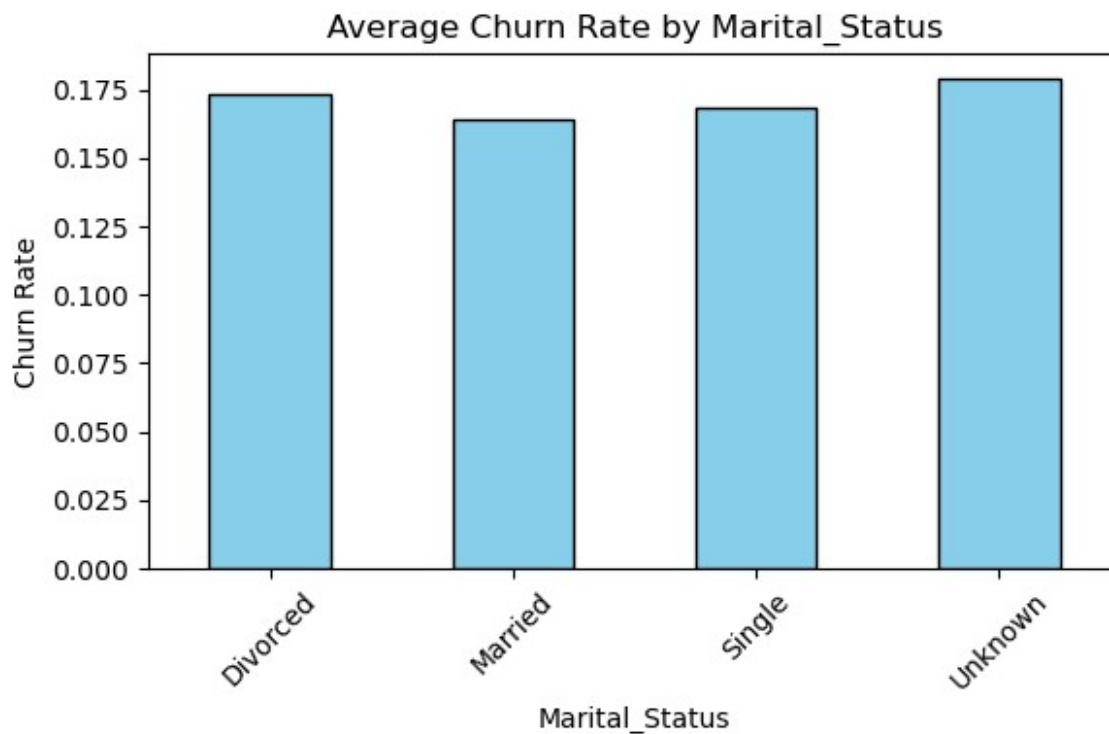
Unknown 0.179012

Name: Churn, dtype: float64

Chi-square statistic: 0.8607

p-value: 0.8349

Degrees of freedom: 3



=== Analysis for Income_Category ===

Churn Rate by Category:

Income_Category

\$120K + 0.181818

\$40K - \$60K 0.158631

\$60K - \$80K 0.140199

\$80K - \$120K 0.168539

Less than \$40K 0.178558

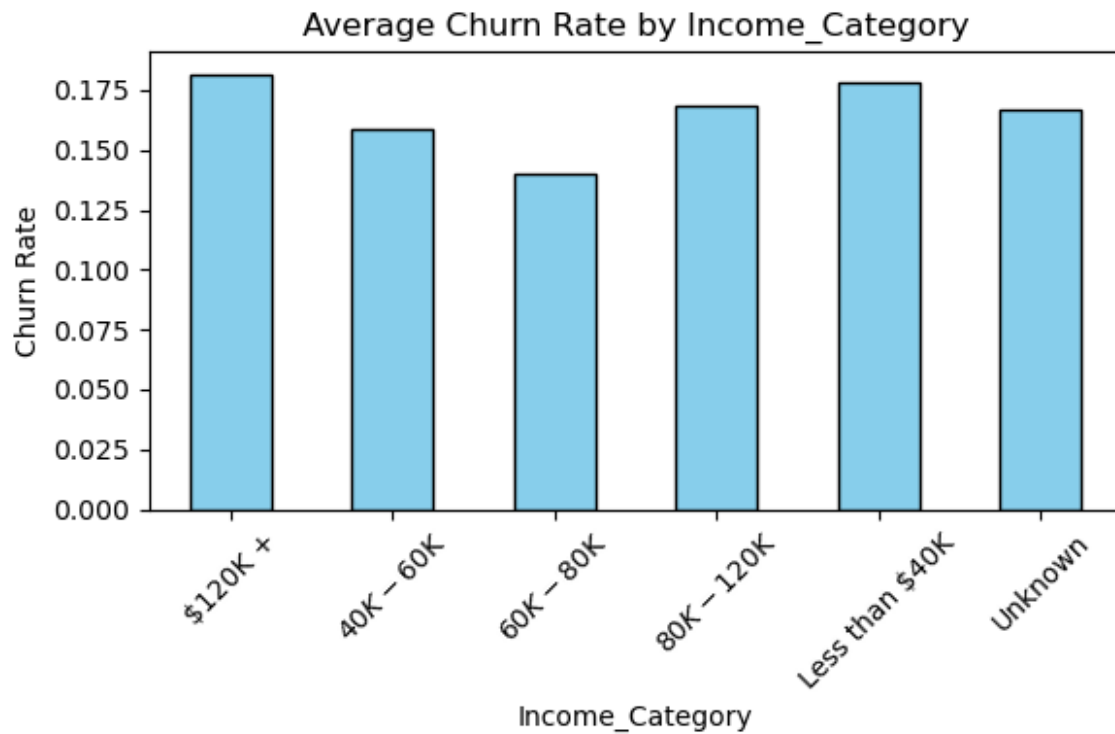
Unknown 0.167131

Name: Churn, dtype: float64

Chi-square statistic: 7.7728

p-value: 0.1692

Degrees of freedom: 5



```
=== Analysis for Card_Category ===
```

```
Churn Rate by Category:
```

```
Card_Category
```

```
Blue      0.167955
```

```
Gold      0.187500
```

```
Platinum  0.666667
```

```
Silver    0.126214
```

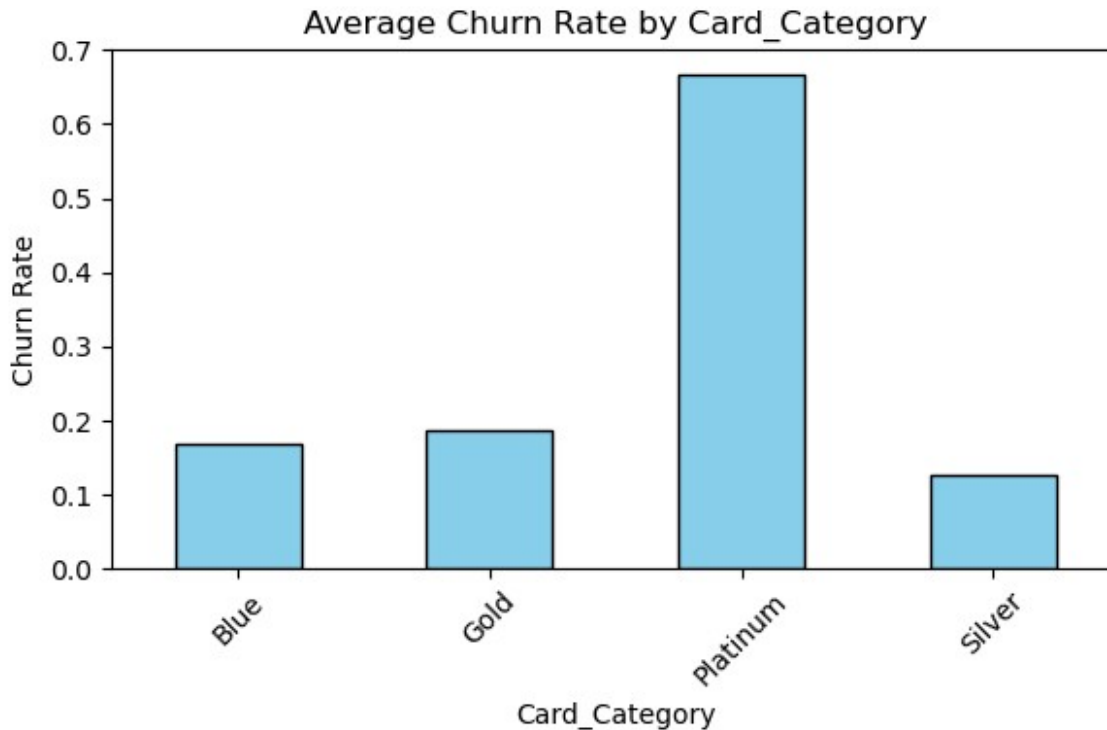
```
Name: Churn, dtype: float64
```

```
Chi-square statistic: 6.6725
```

```
p-value: 0.0831
```

```
Degrees of freedom: 3
```

```
-----
```



4.5 Numerical Feature Distributions

```
import matplotlib.patches as mpatches

numeric_cols = ["Customer_Age", "Credit_Limit", "Total_Trans_Amt"]

# Define custom legend patches
attrited_patch = mpatches.Patch(color='orange', label='Attrited')
existing_patch = mpatches.Patch(color='steelblue', label='Existing')

for col in numeric_cols:
    if col in df.columns:
        plt.figure(figsize=(7, 4))
        sns.histplot(
            data=df,
            x=col,
            hue='Churn',
            kde=True,
            bins=30,
            palette={1: 'orange', 0: 'steelblue'},
            element='step'
        )
        plt.title(f'{col} Distribution by Churn')
        plt.xlabel(col)
        plt.ylabel('Count')

# Custom legend
```

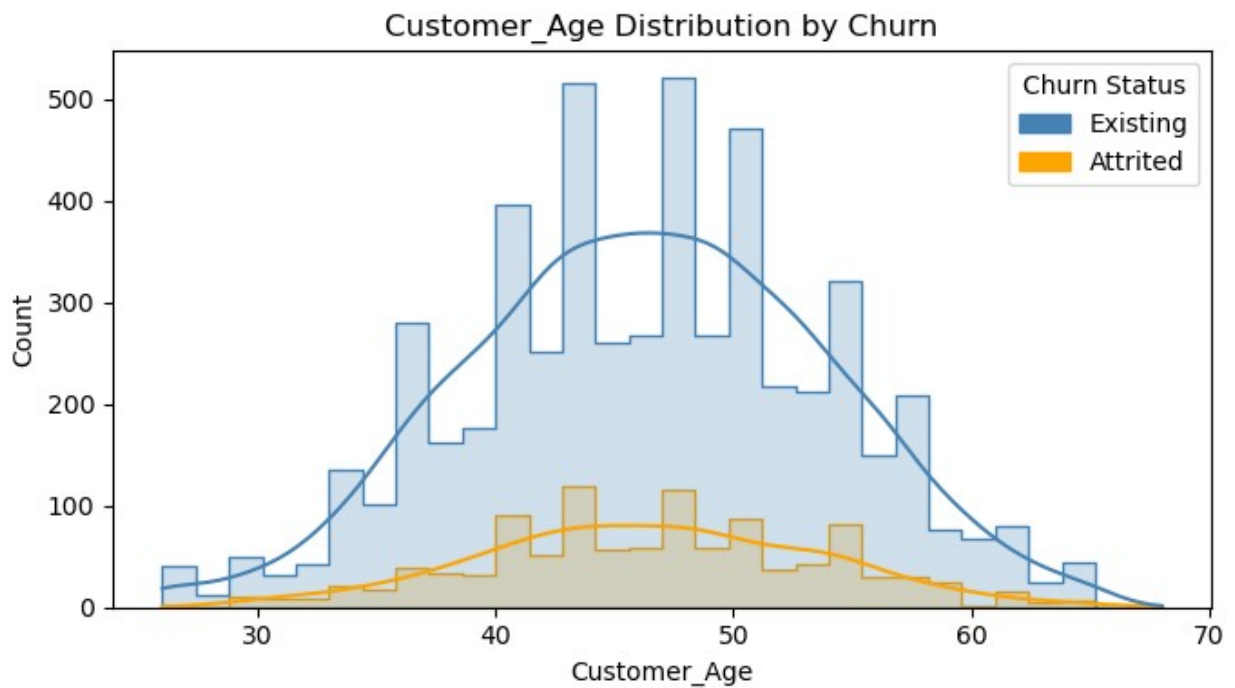


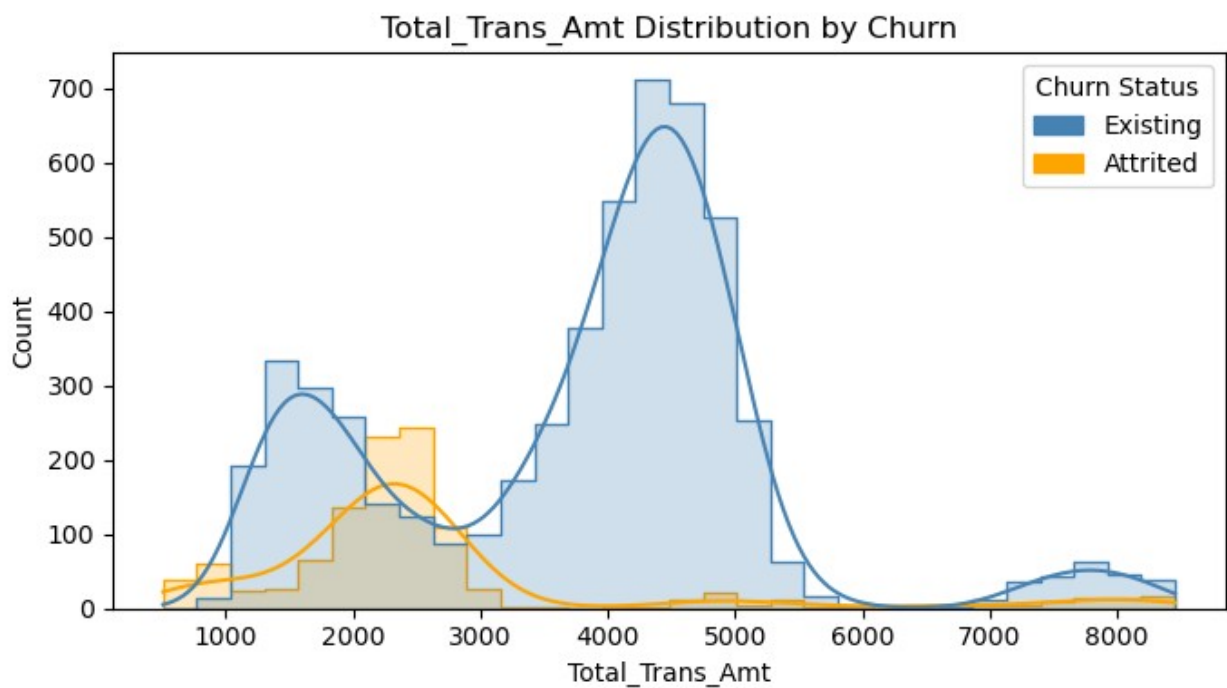
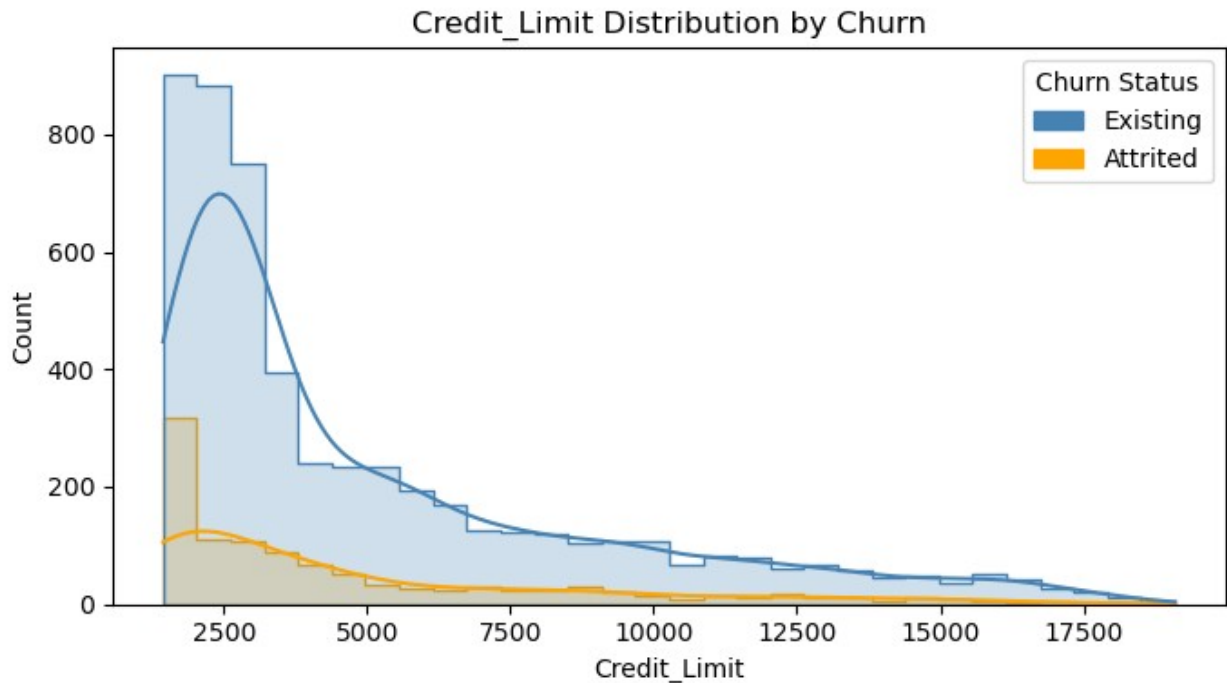
```

plt.legend(
    handles=[existing_patch, attrited_patch],
    title='Churn Status',
    loc='upper right'
)

plt.tight_layout()
plt.show()
else:
    print(f"Column {col} not found.")

```





4.6 Correlation Analysis

We compute the correlation matrix for numeric features and display the values in an annotated heatmap.

```
numeric_features = df.select_dtypes(include=[np.number])
corr_matrix = numeric_features.corr()
```

```

print("Correlation Matrix:")
print(corr_matrix)

plt.figure(figsize=(12, 10))
cax = plt.imshow(corr_matrix, cmap='coolwarm',
interpolation='nearest')
plt.colorbar(cax)
plt.xticks(range(len(corr_matrix.columns)), corr_matrix.columns,
rotation=90)
plt.yticks(range(len(corr_matrix.columns)), corr_matrix.columns)

# Annotate each cell
for i in range(len(corr_matrix.columns)):
    for j in range(len(corr_matrix.columns)):
        value = corr_matrix.iloc[i, j]
        plt.text(j, i, f"{value:.2f}", ha="center", va="center",
color="black")

plt.title('Correlation Matrix', pad=20)
plt.tight_layout()
plt.show()

```

Correlation Matrix:

	Customer_Age	Dependent_count	
Months_on_book \			
Customer_Age	1.000000	-0.135834	
0.749380			
Dependent_count	-0.135834	1.000000	-
0.118447			
Months_on_book	0.749380	-0.118447	
1.000000			
Total_Relationship_Count	-0.027562	-0.053247	-
0.020789			
Months_Inactive_12_mon	0.041709	-0.023846	
0.059521			
Contacts_Count_12_mon	-0.017703	-0.039636	
0.001633			
Credit_Limit	0.001984	0.031626	
0.002032			
Total_Revolving_Bal	0.024946	0.000589	
0.013704			
Avg_Open_To_Buy	-0.003204	0.031405	-
0.000821			
Total_Amt_Chng_Q4_Q1	-0.077916	-0.029982	-
0.054660			
Total_Trans_Amt	-0.004178	0.083532	-
0.011554			
Total_Trans_Ct	-0.045618	0.066685	-
0.040278			

Total_Ct_Chng_Q4_Q1	-0.041602	0.009668	-
0.043758			
Avg_Utilization_Ratio	0.018909	-0.019429	
0.006541			
Churn	0.009278	0.024272	-
0.000709			

	Total_Relationship_Count	
Months_Inactive_12_mon \		
Customer_Age	-0.027562	
0.041709		
Dependent_count	-0.053247	-
0.023846		
Months_on_book	-0.020789	
0.059521		
Total_Relationship_Count	1.000000	-
0.017062		
Months_Inactive_12_mon	-0.017062	
1.000000		
Contacts_Count_12_mon	0.069910	
0.030453		
Credit_Limit	0.024691	-
0.010929		
Total_Revolving_Bal	0.033100	-
0.051495		
Avg_Open_To_Buy	0.017737	-
0.000198		
Total_Amt_Chng_Q4_Q1	0.057729	-
0.015678		
Total_Trans_Amt	-0.214060	-
0.047059		
Total_Trans_Ct	-0.096864	-
0.069923		
Total_Ct_Chng_Q4_Q1	0.024223	-
0.050469		
Avg_Utilization_Ratio	0.028920	-
0.034244		
Churn	-0.166143	
0.200456		

	Contacts_Count_12_mon	Credit_Limit \
Customer_Age	-0.017703	0.001984
Dependent_count	-0.039636	0.031626
Months_on_book	0.001633	0.002032
Total_Relationship_Count	0.069910	0.024691
Months_Inactive_12_mon	0.030453	-0.010929
Contacts_Count_12_mon	1.000000	0.025095
Credit_Limit	0.025095	1.000000
Total_Revolving_Bal	-0.032580	0.089066

Avg_Open_To_Buy	0.031784	0.978361
Total_Amt_Chng_Q4_Q1	-0.001769	-0.027540
Total_Trans_Amt	-0.205141	-0.022577
Total_Trans_Ct	-0.183845	-0.061805
Total_Ct_Chng_Q4_Q1	-0.102562	-0.047357
Avg_Utilization_Ratio	-0.049206	-0.442087
Churn	0.126290	-0.040464

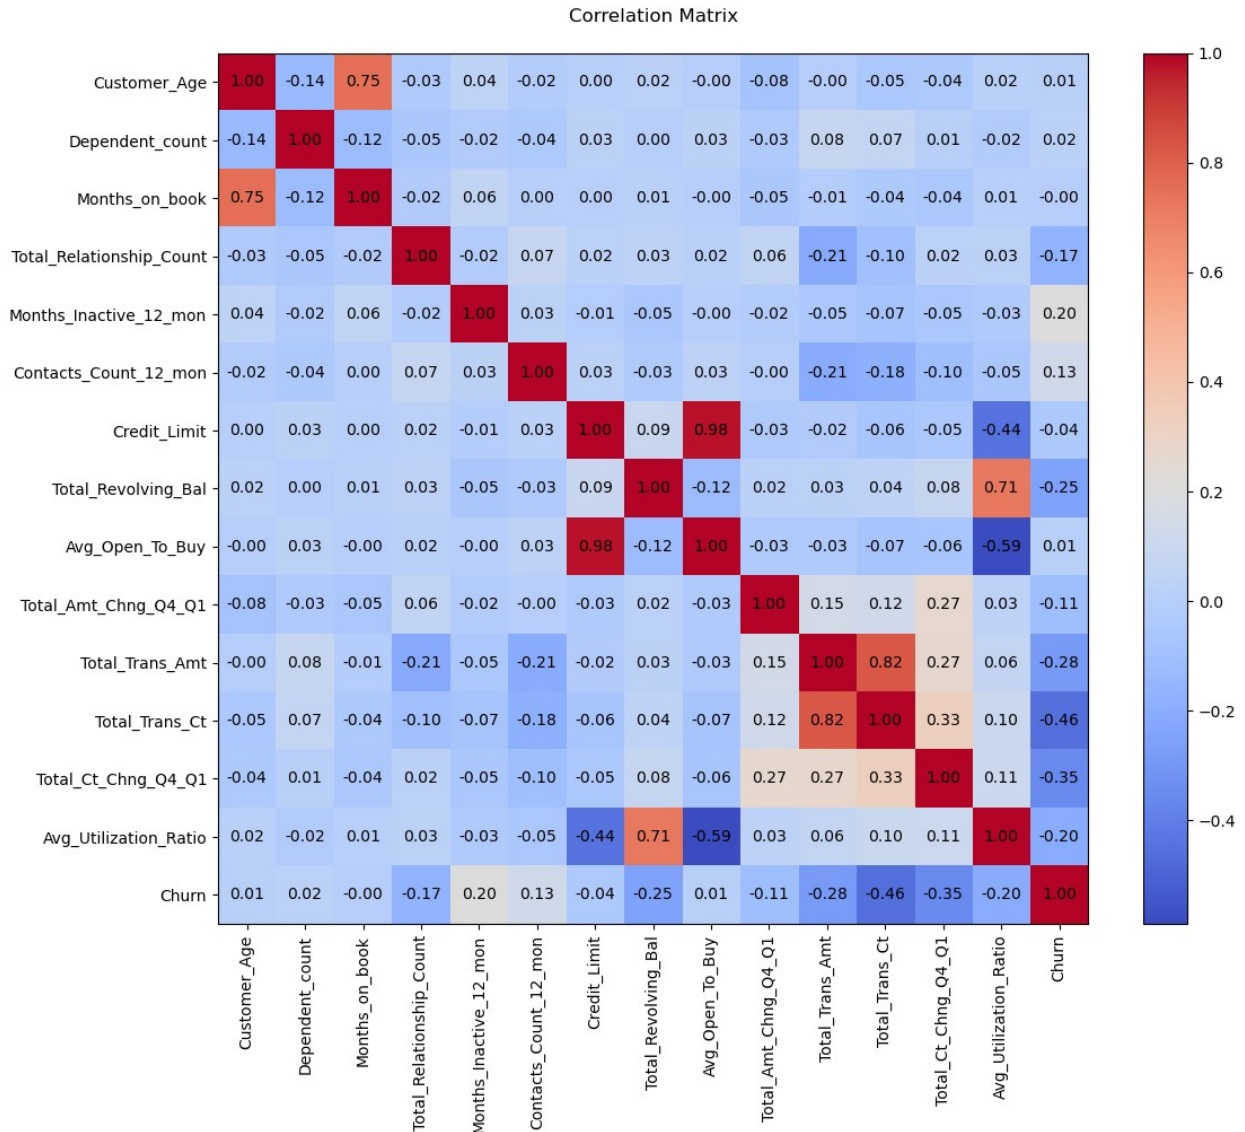
	Total_Revolving_Bal	Avg_Open_To_Buy \
Customer_Age	0.024946	-0.003204
Dependent_count	0.000589	0.031405
Months_on_book	0.013704	-0.000821
Total_Relationship_Count	0.033100	0.017737
Months_Inactive_12_mon	-0.051495	-0.000198
Contacts_Count_12_mon	-0.032580	0.031784
Credit_Limit	0.089066	0.978361
Total_Revolving_Bal	1.000000	-0.118943
Avg_Open_To_Buy	-0.118943	1.000000
Total_Amt_Chng_Q4_Q1	0.019511	-0.031507
Total_Trans_Amt	0.031112	-0.028969
Total_Trans_Ct	0.041654	-0.070264
Total_Ct_Chng_Q4_Q1	0.080912	-0.064016
Avg_Utilization_Ratio	0.714490	-0.589121
Churn	-0.250049	0.011606

	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt \
Customer_Age	-0.077916	-0.004178
Dependent_count	-0.029982	0.083532
Months_on_book	-0.054660	-0.011554
Total_Relationship_Count	0.057729	-0.214060
Months_Inactive_12_mon	-0.015678	-0.047059
Contacts_Count_12_mon	-0.001769	-0.205141
Credit_Limit	-0.027540	-0.022577
Total_Revolving_Bal	0.019511	0.031112
Avg_Open_To_Buy	-0.031507	-0.028969
Total_Amt_Chng_Q4_Q1	1.000000	0.145875
Total_Trans_Amt	0.145875	1.000000
Total_Trans_Ct	0.124872	0.823581
Total_Ct_Chng_Q4_Q1	0.266692	0.270093
Avg_Utilization_Ratio	0.033382	0.062679
Churn	-0.105566	-0.282100

	Total_Trans_Ct	Total_Ct_Chng_Q4_Q1 \
Customer_Age	-0.045618	-0.041602
Dependent_count	0.066685	0.009668
Months_on_book	-0.040278	-0.043758
Total_Relationship_Count	-0.096864	0.024223
Months_Inactive_12_mon	-0.069923	-0.050469
Contacts_Count_12_mon	-0.183845	-0.102562
Credit_Limit	-0.061805	-0.047357

Total_Revolving_Bal	0.041654	0.080912
Avg_Open_To_Buy	-0.070264	-0.064016
Total_Amt_Chng_Q4_Q1	0.124872	0.266692
Total_Trans_Amt	0.823581	0.270093
Total_Trans_Ct	1.000000	0.328686
Total_Ct_Chng_Q4_Q1	0.328686	1.000000
Avg_Utilization_Ratio	0.102378	0.109712
Churn	-0.456916	-0.351080

	Avg_Utilization_Ratio	Churn
Customer_Age	0.018909	0.009278
Dependent_count	-0.019429	0.024272
Months_on_book	0.006541	-0.000709
Total_Relationship_Count	0.028920	-0.166143
Months_Inactive_12_mon	-0.034244	0.200456
Contacts_Count_12_mon	-0.049206	0.126290
Credit_Limit	-0.442087	-0.040464
Total_Revolving_Bal	0.714490	-0.250049
Avg_Open_To_Buy	-0.589121	0.011606
Total_Amt_Chng_Q4_Q1	0.033382	-0.105566
Total_Trans_Amt	0.062679	-0.282100
Total_Trans_Ct	0.102378	-0.456916
Total_Ct_Chng_Q4_Q1	0.109712	-0.351080
Avg_Utilization_Ratio	1.000000	-0.204658
Churn	-0.204658	1.000000



```
import pandas as pd
from sklearn.feature_selection import mutual_info_classif
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_selection import f_classif

# Step 1: Label Encode the categorical columns
label_encoder = LabelEncoder()

# Apply label encoding for each categorical column in the dataset
for column in df.select_dtypes(include=['object']).columns:
    df[column] = label_encoder.fit_transform(df[column])

# Split features and target
```

```

X = df.drop(columns=['Churn']) # Features
y = df['Churn'] # Target

# Step 2: Correlation Coefficient Calculation (Pearson Correlation)
correlation_scores = {}

for feature in X.columns:
    correlation_scores[feature] = X[feature].corr(y) # Pearson correlation

# Step 3: F-Score (ANOVA F-statistic) Calculation
f_scores, p_values = f_classif(X, y)

# Store the F-scores and their corresponding p-values
f_scores_dict = dict(zip(X.columns, f_scores))

# Step 4: Mutual Information Calculation
mutual_info = mutual_info_classif(X, y)

# Store the mutual information results
mi_scores = dict(zip(X.columns, mutual_info))

# Step 5: Rank the results from highest to lowest
sorted_correlation = sorted(correlation_scores.items(), key=lambda x: abs(x[1]), reverse=True)
sorted_f_scores = sorted(f_scores_dict.items(), key=lambda x: x[1], reverse=True)
sorted_mutual_info = sorted(mi_scores.items(), key=lambda x: x[1], reverse=True)

# Step 6: Print the ranked results

# Print Correlation Coefficients ranked from highest to lowest
print("Correlation Coefficients ranked from highest to lowest:")
for feature, corr in sorted_correlation:
    print(f"{feature}: Correlation = {corr:.4f}")

# Print F-scores ranked from highest to lowest
print("\nF-Scores (ANOVA F-statistic) ranked from highest to lowest:")
for feature, f_score in sorted_f_scores:
    print(f"{feature}: F-Score = {f_score:.4f}")

# Print Mutual Information ranked from highest to lowest
print("\nMutual Information ranked from highest to lowest:")
for feature, mi in sorted_mutual_info:
    print(f"{feature}: Mutual Information = {mi:.4f}")

Correlation Coefficients ranked from highest to lowest:
Attrition_Flag: Correlation = -1.0000
Total_Trans_Ct: Correlation = -0.4569
Total_Ct_Chng_Q4_Q1: Correlation = -0.3511

```


Total_Trans_Amt: Correlation = -0.2821
Total_Revolving_Bal: Correlation = -0.2500
Avg_Utilization_Ratio: Correlation = -0.2047
Months_Inactive_12_mon: Correlation = 0.2005
Total_Relationship_Count: Correlation = -0.1661
Contacts_Count_12_mon: Correlation = 0.1263
Total_Amt_Chng_Q4_Q1: Correlation = -0.1056
Credit_Limit: Correlation = -0.0405
Gender: Correlation = -0.0373
Dependent_count: Correlation = 0.0243
Income_Category: Correlation = 0.0177
Avg_Open_To_Buy: Correlation = 0.0116
Card_Category: Correlation = -0.0103
Customer_Age: Correlation = 0.0093
Marital_Status: Correlation = 0.0060
Education_Level: Correlation = 0.0035
Months_on_book: Correlation = -0.0007

F-Scores (ANOVA F-statistic) ranked from highest to lowest:

Attrition_Flag: F-Score = inf
Total_Trans_Ct: F-Score = 1704.7922
Total_Ct_Chng_Q4_Q1: F-Score = 908.3210
Total_Trans_Amt: F-Score = 558.6241
Total_Revolving_Bal: F-Score = 430.9144
Avg_Utilization_Ratio: F-Score = 282.4473
Months_Inactive_12_mon: F-Score = 270.4880
Total_Relationship_Count: F-Score = 183.4083
Contacts_Count_12_mon: F-Score = 104.7175
Total_Amt_Chng_Q4_Q1: F-Score = 72.8146
Credit_Limit: F-Score = 10.5963
Gender: F-Score = 8.9881
Dependent_count: F-Score = 3.8085
Income_Category: F-Score = 2.0217
Avg_Open_To_Buy: F-Score = 0.8703
Card_Category: F-Score = 0.6846
Customer_Age: F-Score = 0.5562
Marital_Status: F-Score = 0.2364
Education_Level: F-Score = 0.0805
Months_on_book: F-Score = 0.0033

Mutual Information ranked from highest to lowest:

Attrition_Flag: Mutual Information = 0.4526
Total_Trans_Amt: Mutual Information = 0.1872
Total_Trans_Ct: Mutual Information = 0.1570
Total_Ct_Chng_Q4_Q1: Mutual Information = 0.1100
Total_Revolving_Bal: Mutual Information = 0.0896
Avg_Utilization_Ratio: Mutual Information = 0.0480
Months_Inactive_12_mon: Mutual Information = 0.0303
Avg_Open_To_Buy: Mutual Information = 0.0211

```
Total_Amt_Chng_Q4_Q1: Mutual Information = 0.0205
Total_Relationship_Count: Mutual Information = 0.0155
Dependent_count: Mutual Information = 0.0076
Contacts_Count_12_mon: Mutual Information = 0.0060
Credit_Limit: Mutual Information = 0.0057
Customer_Age: Mutual Information = 0.0042
Marital_Status: Mutual Information = 0.0042
Months_on_book: Mutual Information = 0.0040
Education_Level: Mutual Information = 0.0017
Income_Category: Mutual Information = 0.0012
Gender: Mutual Information = 0.0000
Card_Category: Mutual Information = 0.0000
```

4.7 Join Plots and Cat Plots

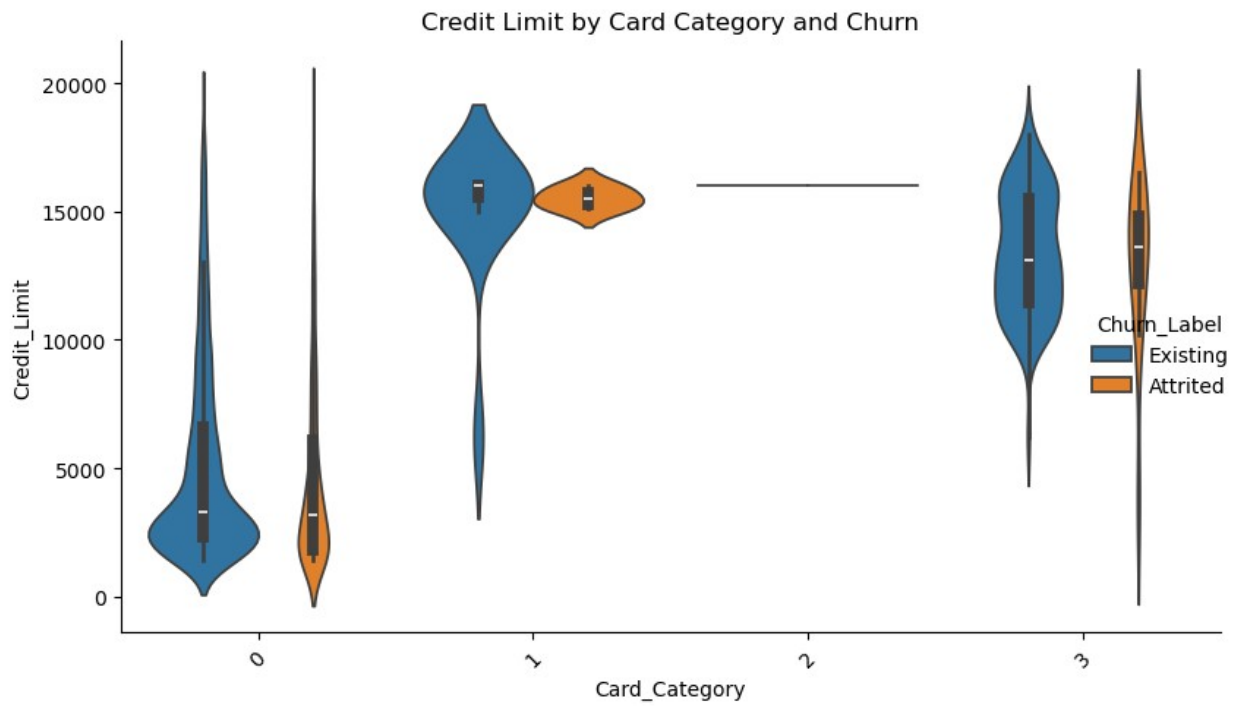
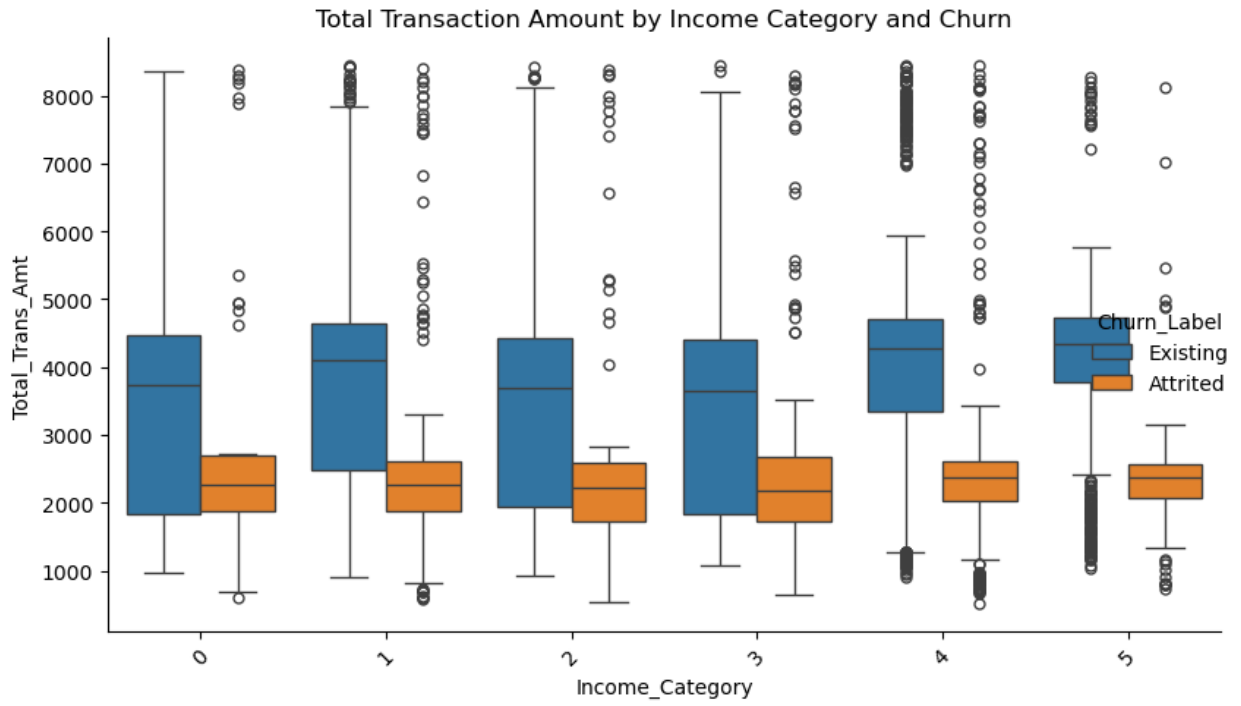
We use Join Plots and Cat Plots to investigate the


```
# Replace 0/1 in 'Churn' with descriptive labels
df['Churn_Label'] = df['Churn'].map({0: 'Existing', 1: 'Attrited'})

# 1. Transaction amount vs Income Category
sns.catplot(data=df, x="Income_Category", y="Total_Trans_Amt",
            hue="Churn_Label", kind="box", height=5, aspect=1.5)
plt.title("Total Transaction Amount by Income Category and Churn")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# 2. Credit Limit by Card Category
sns.catplot(data=df, x="Card_Category", y="Credit_Limit",
            hue="Churn_Label", kind="violin", height=5, aspect=1.5)
plt.title("Credit Limit by Card Category and Churn")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# 3. Count of customers by Education Level and Churn
sns.catplot(data=df, x="Education_Level", hue="Churn_Label",
            kind="count", height=5, aspect=1.5)
plt.title("Churn Count by Education Level")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```







```
# Map churn values for clearer labels
df['Churn_Label'] = df['Churn'].map({0: 'Existing', 1: 'Attrited'})

# 1. Customer Age vs Total Transaction Amount
sns.jointplot(data=df, x="Customer_Age", y="Total_Trans_Amt",
             hue="Churn_Label", kind="scatter", height=6)
plt.suptitle("Customer Age vs Transaction Amount by Churn", y=1.02)
plt.show()

# 2. Credit Limit vs Total Transaction Amount
sns.jointplot(data=df, x="Credit_Limit", y="Total_Trans_Amt",
             hue="Churn_Label", kind="scatter", height=6)
plt.suptitle("Credit Limit vs Transaction Amount by Churn", y=1.02)
plt.show()
```

Customer Age vs Transaction Amount by Churn





5. Predictive Modeling

We build logistic regression, random forest, **XGBoost**, and **AdaBoost** models to predict churn.

5.1 Data Preparation for Modeling

We:

- Drop `Attrition_Flag` from the dataset.
- One-hot encode the known categorical columns.
- Define features (X) and target (y).
- Perform a train-test split.

```

!pip install graphviz
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import chi2_contingency
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier
from sklearn.metrics import (classification_report, confusion_matrix,
                             ConfusionMatrixDisplay, roc_curve,
                             roc_auc_score,
                             precision_recall_fscore_support)

Requirement already satisfied: graphviz in
/opt/homebrew/anaconda3/lib/python3.12/site-packages (0.20.3)

import pandas as pd

# Load the CSV
df = pd.read_csv("BankChurners.csv")

# Drop the 'CLIENTNUM' column
df.drop(columns=['CLIENTNUM'], inplace=True)
df.drop(columns=['Naive_Bayes_Classifier_Attrition_Flag_Card_Category_
Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_
12_mon_1'], inplace=True)
df.drop(columns=['Naive_Bayes_Classifier_Attrition_Flag_Card_Category_
Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_
12_mon_2'], inplace=True)

# Select numeric and categorical columns
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns
categorical_cols = df.select_dtypes(include=['object']).columns

# Print unique values for each categorical column
for col in categorical_cols:
    print(f"Unique values in '{col}':")
    print(df[col].unique())
    print("-" * 50)

Unique values in 'Attrition_Flag':
['Existing Customer' 'Attrited Customer']
-----

Unique values in 'Gender':
['M' 'F']
-----

```

```

Unique values in 'Education_Level':
['High School' 'Graduate' 'Uneducated' 'Unknown' 'College' 'Post-
Graduate'
'Doctorate']
-----
Unique values in 'Marital_Status':
['Married' 'Single' 'Unknown' 'Divorced']
-----
Unique values in 'Income_Category':
['$60K - $80K' 'Less than $40K' '$80K - $120K' '$40K - $60K' '$120K +'
'Unknown']
-----
Unique values in 'Card_Category':
['Blue' 'Gold' 'Silver' 'Platinum']
-----

```

```

cleaned_df = df.copy()
# Drop NA
cleaned_df.dropna()
# Drop duplicates
cleaned_df.drop_duplicates(inplace=True)

# IQR Filtering
for col in numeric_cols:
    Q1 = cleaned_df[col].quantile(0.25)
    Q3 = cleaned_df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    IQR_cleaned_df = cleaned_df[(cleaned_df[col] >= lower_bound) &
(cleaned_df[col] <= upper_bound)]

# Step 4: Check final shape after removing outliers
print("Cleaned dataset shape:", cleaned_df.shape)

# Optional: Save cleaned data
cleaned_df.to_csv("IQRcleaned_dataset.csv", index=False)
cleaned_df.head

```

Cleaned dataset shape: (10127, 20)

<bound	method	NDFrame.head of		Attrition_Flag	Customer_Age
Gender	Dependent_count	\			
0	Existing Customer	45	M		3
1	Existing Customer	49	F		5
2	Existing Customer	51	M		3
3	Existing Customer	40	F		4
4	Existing Customer	40	M		3
...
10122	Existing Customer	50	M		2

10123	Attrited Customer	41	M	2
10124	Attrited Customer	44	F	1
10125	Attrited Customer	30	M	2
10126	Attrited Customer	43	F	2

	Education_Level	Marital_Status	Income_Category	Card_Category \
0	High School	Married	\$60K - \$80K	Blue
1	Graduate	Single	Less than \$40K	Blue
2	Graduate	Married	\$80K - \$120K	Blue
3	High School	Unknown	Less than \$40K	Blue
4	Uneducated	Married	\$60K - \$80K	Blue
...
10122	Graduate	Single	\$40K - \$60K	Blue
10123	Unknown	Divorced	\$40K - \$60K	Blue
10124	High School	Married	Less than \$40K	Blue
10125	Graduate	Unknown	\$40K - \$60K	Blue
10126	Graduate	Married	Less than \$40K	Silver

	Months_on_book	Total_Relationship_Count
Months_Inactive_12_mon \		
0	39	5
1		
1	44	6
1		
2	36	4
1		
3	34	3
4		
4	21	5
1		
...
.		
10122	40	3
2		
10123	25	4
2		
10124	36	5
3		
10125	36	4
3		
10126	25	6
2		

	Contacts_Count_12_mon	Credit_Limit	Total_Revolving_Bal \
0	3	12691.0	777
1	2	8256.0	864
2	0	3418.0	0
3	1	3313.0	2517
4	0	4716.0	0
...

10122	3	4003.0	1851
10123	3	4277.0	2186
10124	4	5409.0	0
10125	3	5281.0	0
10126	4	10388.0	1961

	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt
Total_Trans_Ct \			
0	11914.0	1.335	1144
42			
1	7392.0	1.541	1291
33			
2	3418.0	2.594	1887
20			
3	796.0	1.405	1171
20			
4	4716.0	2.175	816
28			
...
...			
10122	2152.0	0.703	15476
117			
10123	2091.0	0.804	8764
69			
10124	5409.0	0.819	10291
60			
10125	5281.0	0.535	8395
62			
10126	8427.0	0.703	10294
61			

	Total_Ct_Chng_Q4_Q1	Avg_Utilization_Ratio
0	1.625	0.061
1	3.714	0.105
2	2.333	0.000
3	2.333	0.760
4	2.500	0.000
...
10122	0.857	0.462
10123	0.683	0.511
10124	0.818	0.000
10125	0.722	0.000
10126	0.649	0.189

[10127 rows x 20 columns]>

5.2 Finding the Most Relevant Features

```
import pandas as pd
import seaborn as sns
```

```

import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder

# List of categorical columns as provided.
categorical_cols = [
    'Attrition_Flag',
    'Gender',
    'Education_Level',
    'Marital_Status',
    'Income_Category',
    'Card_Category'
]

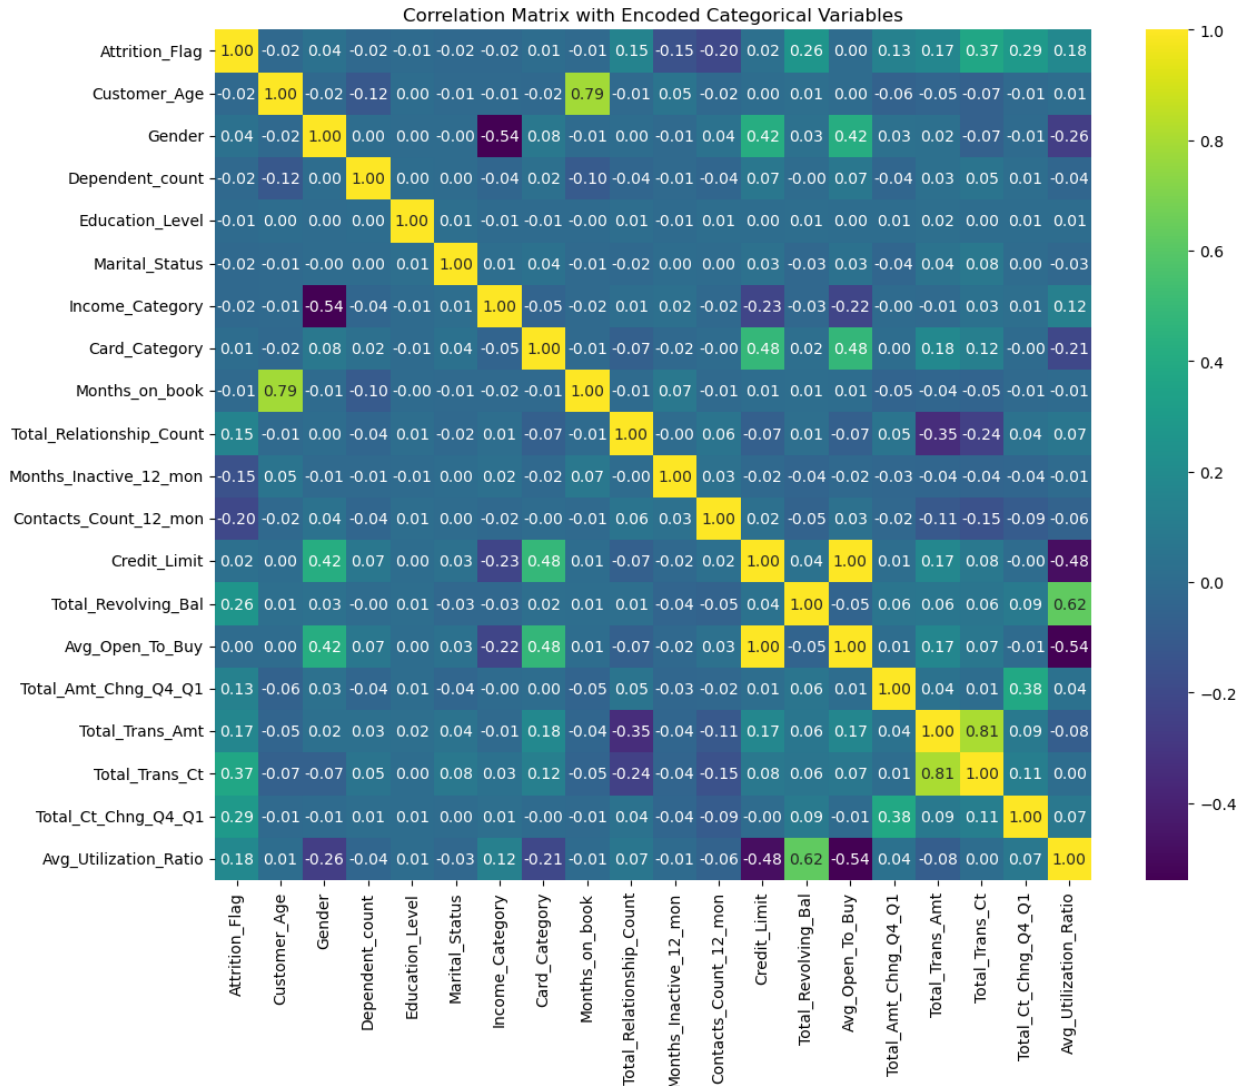
# Create a copy of the DataFrame for encoding.
df_encoded = IQR_cleaned_df.copy()

# Apply label encoding to each categorical column.
le = LabelEncoder()
for col in categorical_cols:
    df_encoded[col] = le.fit_transform(df_encoded[col])

# Compute the correlation matrix for the entire DataFrame.
corr_matrix = df_encoded.corr()

# Create a heatmap for the correlation matrix.
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, cmap='viridis', annot=True, fmt=".2f")
plt.title('Correlation Matrix with Encoded Categorical Variables')
plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()

```



5.2.2 F-Score

```
import pandas as pd
from sklearn.feature_selection import f_classif

# Assume df_encoded is your DataFrame with all features encoded,
# including "Attrition_Flag".
# Define the target variable.
target = 'Attrition_Flag'

# Create X containing all features except the target variable,
# and y containing the target.
X = df_encoded.drop(columns=[target])
y = df_encoded[target]

# Compute F-scores and p-values using the ANOVA F-test.
F_values, p_values = f_classif(X, y)
```

```
# Create a DataFrame to display the F-scores and p-values along with
the feature names.
f_scores_df = pd.DataFrame({
    'Feature': X.columns,
    'F_score': F_values,
    'p_value': p_values
}).sort_values(by='F_score', ascending=False)

print("F-scores against Attrition_Flag:")
print(f_scores_df)
```

```
F-scores against Attrition_Flag:
```

	Feature	F_score	p_value
16	Total_Trans_Ct	1620.121692	0.000000e+00
17	Total_Ct_Chng_Q4_Q1	930.078416	1.647725e-195
12	Total_Revolving_Bal	752.702408	6.630148e-160
10	Contacts_Count_12_mon	441.868050	4.697490e-96
18	Avg_Utilization_Ratio	332.876795	3.357689e-73
15	Total_Trans_Amt	296.227714	1.857439e-65
9	Months_Inactive_12_mon	240.910376	1.032664e-53
8	Total_Relationship_Count	233.072886	4.829281e-52
14	Total_Amt_Chng_Q4_Q1	176.961638	4.836643e-40
1	Gender	14.085007	1.757076e-04
11	Credit_Limit	5.773729	1.628536e-02
2	Dependent_count	3.652825	5.600239e-02
4	Marital_Status	3.503085	6.128344e-02
0	Customer_Age	3.356074	6.698689e-02
5	Income_Category	3.131546	7.682098e-02
7	Months_on_book	1.897071	1.684370e-01
6	Card_Category	0.369161	5.434755e-01
3	Education_Level	0.312007	5.764636e-01
13	Avg_Open_To_Buy	0.000823	9.771161e-01

5.2.3 Mutual Information

```
import pandas as pd
from sklearn.feature_selection import mutual_info_classif

# Assume df_encoded is your DataFrame with features encoded, including
"Attrition_Flag".

# Define the target variable and feature set.
target = 'Attrition_Flag'
X = df_encoded.drop(columns=[target])
y = df_encoded[target]

# Compute the mutual information for each feature with respect to the
target.
mi_values = mutual_info_classif(X, y, random_state=42)
```

```
# Create a DataFrame to display the mutual information values along
with the feature names.
mi_df = pd.DataFrame({
    'Feature': X.columns,
    'Mutual_Information': mi_values
}).sort_values(by='Mutual_Information', ascending=False)

print("Mutual Information of each feature with Attrition_Flag:")
print(mi_df)
```

Mutual Information of each feature with Attrition_Flag:

	Feature	Mutual_Information
15	Total_Trans_Amt	0.159579
16	Total_Trans_Ct	0.113819
17	Total_Ct_Chng_Q4_Q1	0.097490
12	Total_Revolving_Bal	0.081312
18	Avg_Utilization_Ratio	0.046084
14	Total_Amt_Chng_Q4_Q1	0.028307
10	Contacts_Count_12_mon	0.023935
13	Avg_Open_To_Buy	0.020645
9	Months_Inactive_12_mon	0.017864
8	Total_Relationship_Count	0.013571
11	Credit_Limit	0.009097
3	Education_Level	0.006701
7	Months_on_book	0.005862
4	Marital_Status	0.005622
2	Dependent_count	0.005081
0	Customer_Age	0.002752
1	Gender	0.002016
6	Card_Category	0.001135
5	Income_Category	0.000975

5.2.4 Elbow Plot to Determine Minimum Columns Required

```
import pandas as pd
import numpy as np
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier # or
LogisticRegression
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt

# --- Setup ---
# Your DataFrame with encoded categorical columns and target column
df = df_encoded.copy()
y = df['Attrition_Flag']

# Ranked features from previous combined ranking
ranked_features = [
    'Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1', 'Total_Revolving_Bal',
```

```

'Total_Trans_Amt',
  'Avg_Utilization_Ratio', 'Contacts_Count_12_mon',
'Months_Inactive_12_mon',
  'Total_Relationship_Count', 'Total_Amt_Chng_Q4_Q1',
'Credit_Limit', 'Gender',
  'Dependent_count', 'Customer_Age', 'Marital_Status',
'Income_Category',
  'Avg_Open_To_Buy', 'Months_on_book', 'Education_Level',
'Card_Category'
]

# --- Model and scoring ---
model = RandomForestClassifier(random_state=42)
scoring_metric = 'f1' # or 'accuracy'

# --- Elbow plot prep ---
scores = []
num_features_list = list(range(1, len(ranked_features) + 1))

for i in num_features_list:
    selected_features = ranked_features[:i]
    X = df[selected_features]

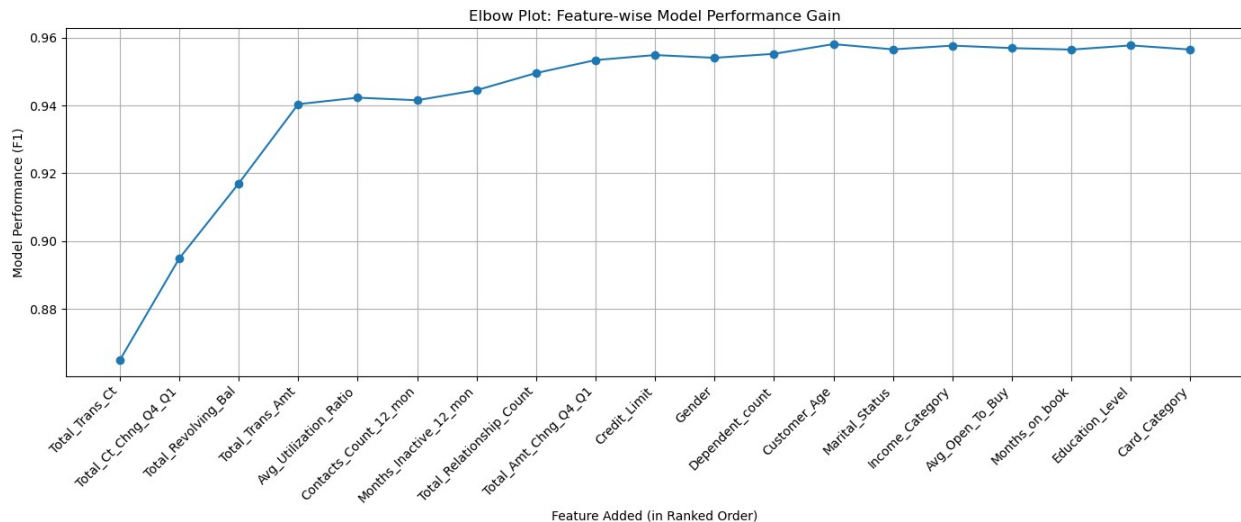
    # 5-fold cross-validation
    score = cross_val_score(model, X, y, cv=5,
scoring=scoring_metric).mean()
    scores.append(score)

# --- Plot ---
# --- Plot ---
plt.figure(figsize=(14, 6))
plt.plot(num_features_list, scores, marker='o')

# Set x-axis labels to the feature name added at each step
added_features = ranked_features # Already ordered by importance
plt.xticks(ticks=num_features_list, labels=added_features,
rotation=45, ha='right')

plt.title('Elbow Plot: Feature-wise Model Performance Gain')
plt.xlabel('Feature Added (in Ranked Order)')
plt.ylabel(f'Model Performance ({scoring_metric.upper()})')
plt.grid(True)
plt.tight_layout()
plt.show()

```



6 Machine Learning Models

6.1 Logistic Regression

Logistic regression estimates the probability of an event occurring, such as voted or didn't vote, based on a given data set of independent variables.

Logistic Regression: Attrition_Flag against Total_Trans_Ct

```
import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Define X and y
X = df_encoded[['Total_Trans_Ct']]
y = df_encoded['Attrition_Flag']

# Fit logistic regression
logreg = LogisticRegression()
logreg.fit(X, y)

# Predict class and probabilities
y_pred_class = logreg.predict(X)
y_pred_proba = logreg.predict_proba(X)[:, 1] # Probability of class 1
(Attrited)

# Evaluation metrics
acc = accuracy_score(y, y_pred_class)
auc = roc_auc_score(y, y_pred_proba)
cm = confusion_matrix(y, y_pred_class)
```



```


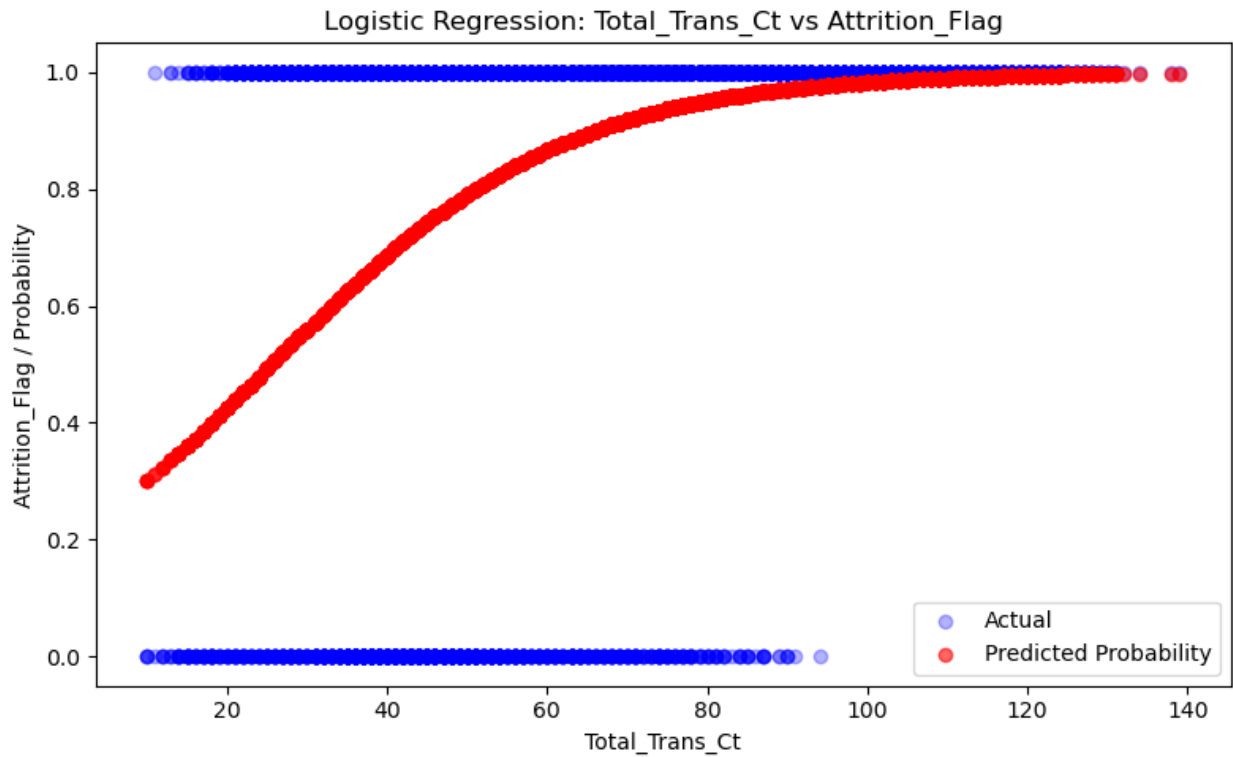
print("Logistic Regression Results:")
print(f"Coefficient: {logreg.coef_[0][0]}")
print(f"Intercept: {logreg.intercept_[0]}")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc:.4f}")
print("Confusion Matrix:")
print(cm)

# --- Visual: Logistic Regression Curve ---
plt.figure(figsize=(8, 5))
plt.scatter(X, y, color='blue', alpha=0.3, label='Actual')
plt.scatter(X, y_pred_proba, color='red', alpha=0.6, label='Predicted Probability')
plt.xlabel('Total_Trans_Ct')
plt.ylabel('Attrition_Flag / Probability')
plt.title('Logistic Regression: Total_Trans_Ct vs Attrition_Flag')
plt.legend()
plt.tight_layout()
plt.show()

# --- Visual: Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Pred: Attrited', 'Pred: Existing'],
            yticklabels=['Actual: Attrited', 'Actual: Existing'])
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()

Logistic Regression Results:
Coefficient: 0.05432217053505641
Intercept: -1.3909493448085435
Accuracy: 0.8334
ROC AUC Score: 0.7956
Confusion Matrix:
[[ 134 1493]
 [ 194 8306]]

```



Logistic Regression: Attrition_Flag against Total_Ct_Chng_Q4_Q1

```
import pandas as pd
import numpy as np
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Define X and y
X = df_encoded[['Total_Ct_Chng_Q4_Q1']]
y = df_encoded['Attrition_Flag']

# Fit logistic regression
logreg = LogisticRegression()
logreg.fit(X, y)

# Predict class and probabilities
y_pred_class = logreg.predict(X)
y_pred_proba = logreg.predict_proba(X)[:, 1] # Probability of class 1
(Attrited)

# Evaluation metrics
acc = accuracy_score(y, y_pred_class)
auc = roc_auc_score(y, y_pred_proba)
cm = confusion_matrix(y, y_pred_class)

print("Logistic Regression Results:")
print(f"Coefficient: {logreg.coef_[0][0]}")
print(f"Intercept: {logreg.intercept_[0]}")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc:.4f}")
print("Confusion Matrix:")
print(cm)

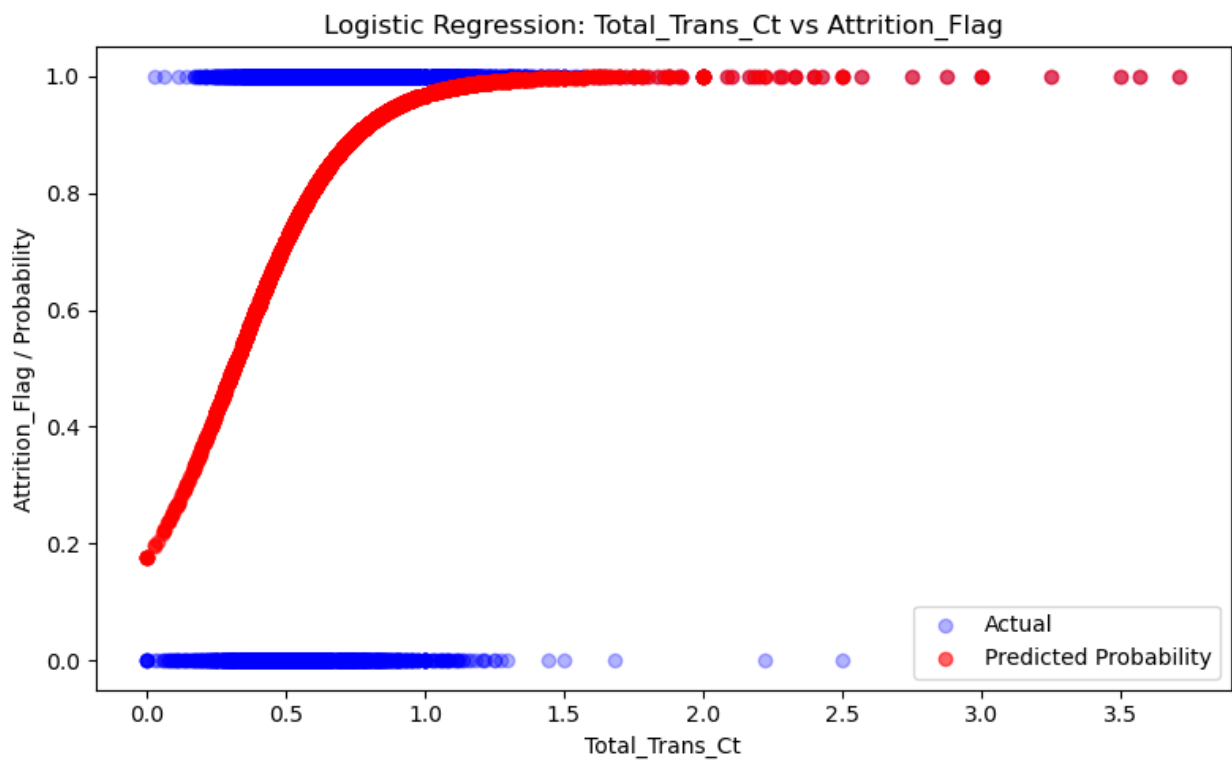
# --- Visual: Logistic Regression Curve ---
plt.figure(figsize=(8, 5))
plt.scatter(X, y, color='blue', alpha=0.3, label='Actual')
plt.scatter(X, y_pred_proba, color='red', alpha=0.6, label='Predicted
Probability')
plt.xlabel('Total_Trans_Ct')
plt.ylabel('Attrition_Flag / Probability')
plt.title('Logistic Regression: Total_Trans_Ct vs Attrition_Flag')
plt.legend()
plt.tight_layout()
plt.show()


# --- Visual: Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Pred: Attrited', 'Pred: Existing'],
            yticklabels=['Actual: Attrited', 'Actual: Existing'])
plt.title('Confusion Matrix Heatmap')

```

```
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()
```

Logistic Regression Results:
Coefficient: 4.965997692835299
Intercept: -1.5473791155267234
Accuracy: 0.8490
ROC AUC Score: 0.7453
Confusion Matrix:
[[187 1440]
 [89 8411]]





Logistic Regression: Attrition_Flag against Total_Revolving_Bal

```
import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Define X and y
X = df_encoded[['Total_Revolving_Bal']]
y = df_encoded['Attrition_Flag']

# Fit logistic regression
logreg = LogisticRegression()
logreg.fit(X, y)

# Predict class and probabilities
y_pred_class = logreg.predict(X)
y_pred_proba = logreg.predict_proba(X)[:, 1] # Probability of class 1
(Attrited)

# Evaluation metrics
acc = accuracy_score(y, y_pred_class)
auc = roc_auc_score(y, y_pred_proba)
cm = confusion_matrix(y, y_pred_class)
```

```



print("Logistic Regression Results:")
print(f"Coefficient: {logreg.coef_[0][0]}")
print(f"Intercept: {logreg.intercept_[0]}")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc:.4f}")
print("Confusion Matrix:")
print(cm)

# --- Visual: Logistic Regression Curve ---
plt.figure(figsize=(8, 5))
plt.scatter(X, y, color='blue', alpha=0.3, label='Actual')
plt.scatter(X, y_pred_proba, color='red', alpha=0.6, label='Predicted Probability')
plt.xlabel('Total_Trans_Ct')
plt.ylabel('Attrition_Flag / Probability')
plt.title('Logistic Regression: Total_Trans_Ct vs Attrition_Flag')
plt.legend()
plt.tight_layout()
plt.show()

# --- Visual: Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Pred: Attrited Customer', 'Pred: Existing Customer'],
            yticklabels=['Actual: Attrited Customer', 'Actual: Existing Customer'])
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()

Logistic Regression Results:
Coefficient: 0.0009275749206390263
Intercept: 0.7633947274499862
Accuracy: 0.8393
ROC AUC Score: 0.6877
Confusion Matrix:
[[ 0 1627]
 [ 0 8500]]

```



Logistic Regression: Attrition_Flag against Total_Trans_Amt

```
import pandas as pd
import numpy as np
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Define X and y
X = df_encoded[['Total_Trans_Amt']]
y = df_encoded['Attrition_Flag']

# Fit logistic regression
logreg = LogisticRegression()
logreg.fit(X, y)

# Predict class and probabilities
y_pred_class = logreg.predict(X)
y_pred_proba = logreg.predict_proba(X)[:, 1] # Probability of class 1
(Attrited)

# Evaluation metrics
acc = accuracy_score(y, y_pred_class)
auc = roc_auc_score(y, y_pred_proba)
cm = confusion_matrix(y, y_pred_class)

print("Logistic Regression Results:")
print(f"Coefficient: {logreg.coef_[0][0]}")
print(f"Intercept: {logreg.intercept_[0]}")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc:.4f}")
print("Confusion Matrix:")
print(cm)

# --- Visual: Logistic Regression Curve ---
plt.figure(figsize=(8, 5))
plt.scatter(X, y, color='blue', alpha=0.3, label='Actual')
plt.scatter(X, y_pred_proba, color='red', alpha=0.6, label='Predicted
Probability')
plt.xlabel('Total_Trans_Ct')
plt.ylabel('Attrition_Flag / Probability')
plt.title('Logistic Regression: Total_Trans_Ct vs Attrition_Flag')
plt.legend()
plt.tight_layout()
plt.show()

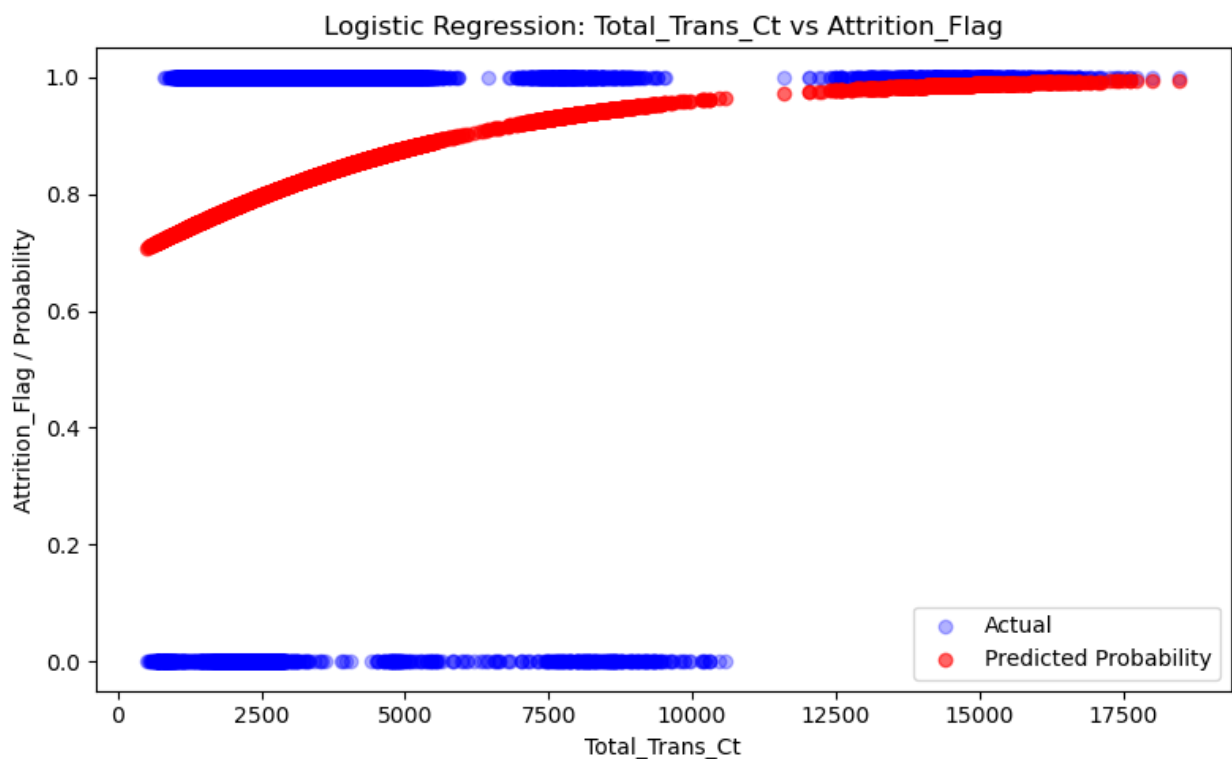
# --- Visual: Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Pred: Attrited Customer', 'Pred: Existing
Customer'],
            yticklabels=['Actual: Attrited Customer', 'Actual:


```



```
Existing Customer'])
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()
```

Logistic Regression Results:
Coefficient: 0.00024092367759846606
Intercept: 0.7639090174091386
Accuracy: 0.8393
ROC AUC Score: 0.6759
Confusion Matrix:
[[0 1627]
[0 8500]]





Logistic Regression: Attrition_Flag against df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1', 'Total_Revolving_Bal', 'Total_Trans_Amt']]

In view of the low accuracy and performance of single variables, we decided to use a multivariate feature to predict Attrition_Flag

```
import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Define X and y with 4 top-ranked features
X = df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
'Total_Revolving_Bal', 'Total_Trans_Amt']]
y = df_encoded['Attrition_Flag']

# Fit logistic regression
logreg = LogisticRegression(max_iter=1000) # Increase max_iter for
convergence
logreg.fit(X, y)

# Predict class and probabilities
y_pred_class = logreg.predict(X)
y_pred_proba = logreg.predict_proba(X)[:, 1] # Probability of class 1
```

(Attrited)

Evaluation metrics

```
acc = accuracy_score(y, y_pred_class)
auc = roc_auc_score(y, y_pred_proba)
cm = confusion_matrix(y, y_pred_class)
```

```
print("Logistic Regression Results:")
for i, feature in enumerate(X.columns):
    print(f"Coefficient for {feature}: {logreg.coef_[0][i]:.4f}")
print(f"Intercept: {logreg.intercept_[0]:.4f}")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc:.4f}")
print("Confusion Matrix:")
print(cm)
```

--- Visual: Confusion Matrix Heatmap ---

```
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Pred: Attrited', 'Pred: Existing'],
            yticklabels=['Actual: Attrited', 'Actual: Existing'])
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()
```

Compute linear combination (logit) manually


```
logit_values = np.dot(X, logreg.coef_[0]) + logreg.intercept_[0]
```

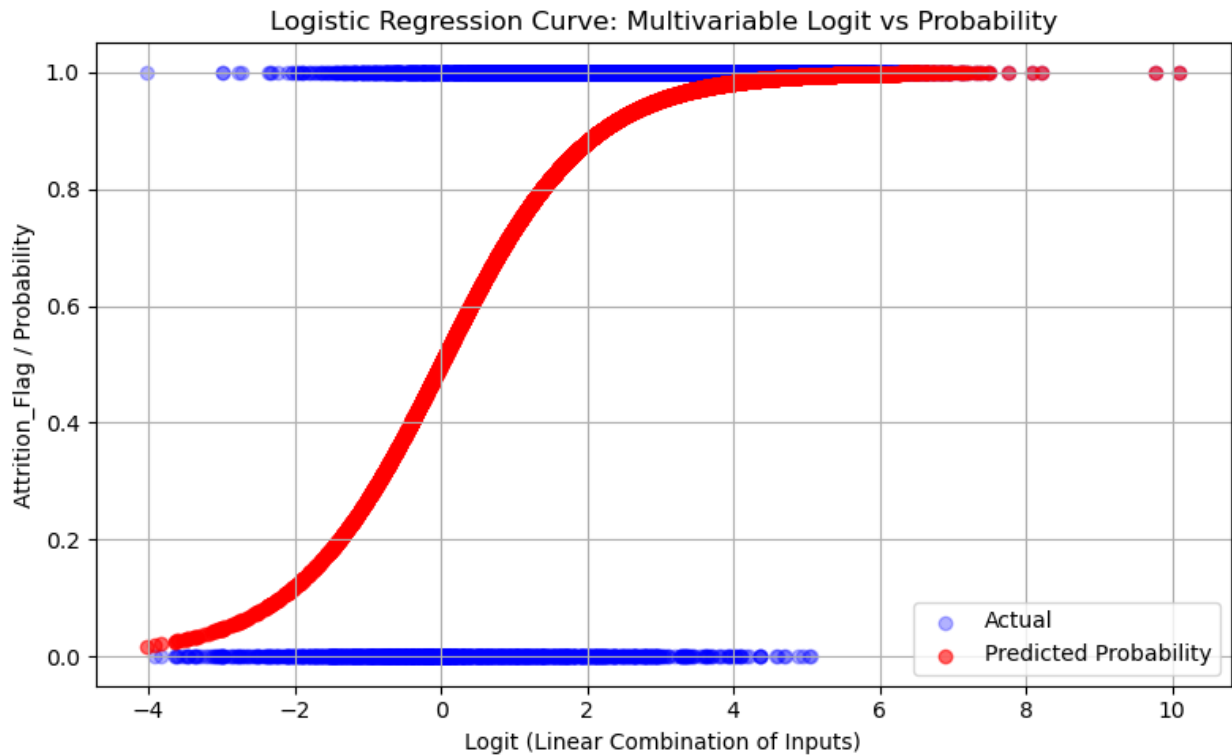
Plot the sigmoid curve

```
plt.figure(figsize=(8, 5))
plt.scatter(logit_values, y, color='blue', alpha=0.3, label='Actual')
plt.scatter(logit_values, y_pred_proba, color='red', alpha=0.6,
            label='Predicted Probability')
plt.xlabel('Logit (Linear Combination of Inputs)')
plt.ylabel('Attrition_Flag / Probability')
plt.title('Logistic Regression Curve: Multivariable Logit vs Probability')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
Logistic Regression Results:
Coefficient for Total_Trans_Ct: 0.1040
Coefficient for Total_Ct_Chng_Q4_Q1: 3.0909
Coefficient for Total_Revolving_Bal: 0.0010
Coefficient for Total_Trans_Amt: -0.0005
Intercept: -5.4061
```

Accuracy: 0.8851
ROC AUC Score: 0.8938
Confusion Matrix:
[[805 822]
 [342 8158]]





6.2 Decision Trees

```
import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix, roc_curve, auc
import matplotlib.pyplot as plt
import seaborn as sns

# Define X and y with top 4 features
X = df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
'Total_Revolving_Bal', 'Total_Trans_Amt']]
y = df_encoded['Attrition_Flag']

# Fit Decision Tree Classifier
tree_clf = DecisionTreeClassifier(max_depth=4, random_state=42) # You
can tweak max_depth
tree_clf.fit(X, y)

# Predictions
y_pred_class = tree_clf.predict(X)
y_pred_proba = tree_clf.predict_proba(X)[:, 1]

# Evaluation metrics
acc = accuracy_score(y, y_pred_class)
auc_score = roc_auc_score(y, y_pred_proba)
```

```

cm = confusion_matrix(y, y_pred_class)

print("Decision Tree Classifier Results:")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc_score:.4f}")
print("Confusion Matrix:")
print(cm)

# --- Visual: Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Greens', cbar=False,
            xticklabels=['Pred: Attrited', 'Pred: Existing'],
            yticklabels=['Actual: Attrited', 'Actual: Existing'])
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()

# --- Visual: ROC Curve ---
fpr, tpr, _ = roc_curve(y, y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Decision Tree ROC Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()



# --- Visual: Decision Tree Structure ---
plt.figure(figsize=(16, 8))
plot_tree(tree_clf, feature_names=X.columns, class_names=['Not Attrited', 'Attrited'],
          filled=True, rounded=True)
plt.title("Decision Tree Structure (max_depth=4)")
plt.tight_layout()
plt.show()

```

```

Decision Tree Classifier Results:
Accuracy: 0.9275
ROC AUC Score: 0.9476
Confusion Matrix:
[[1247  380]
 [ 354 8146]]

```



The diagram illustrates a decision tree for the 'Species' variable. The root node is 'Tree: Species (0.75)'. It splits on 'Species' (0.75) into two main branches: 'Species = 0.75' and 'Species = 1.00'. Each branch further splits on 'Species' (0.75) into 'Species = 0.75' and 'Species = 1.00'. The final leaf nodes represent the predicted species for each combination of splits.

6.3 Bagging of Trees: Random Forest

Random forest uses multiple decision trees to make predictions, improving accuracy and reducing overfitting.

```
import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix, roc_curve, auc
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt
import seaborn as sns

# Define X and y
X = df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
'Total_Revolving_Bal', 'Total_Trans_Amt']]
y = df_encoded['Attrition_Flag']

# Fit Random Forest Classifier
rf_clf = RandomForestClassifier(n_estimators=100, max_depth=5,
random_state=42)
rf_clf.fit(X, y)

# Predictions
y_pred_class = rf_clf.predict(X)
y_pred_proba = rf_clf.predict_proba(X)[:, 1]

# Evaluation metrics
acc = accuracy_score(y, y_pred_class)
auc_score = roc_auc_score(y, y_pred_proba)
```



```

cm = confusion_matrix(y, y_pred_class)

print("Random Forest Classifier Results:")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc_score:.4f}")
print("Confusion Matrix:")
print(cm)

# --- Visual: Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Greens', cbar=False,
            xticklabels=['Pred: Attrited', 'Pred: Existing'],
            yticklabels=['Actual: Attrited', 'Actual: Existing'])
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()

# --- Visual: ROC Curve ---
fpr, tpr, _ = roc_curve(y, y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Random Forest ROC Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Visual: Feature Importances ---
importances = rf_clf.feature_importances_
features = X.columns
indices = np.argsort(importances)[::-1]

plt.figure(figsize=(160, 80), dpi=200)
sns.barplot(x=importances[indices], y=features[indices],
           palette='magma')
plt.title("Feature Importances (Random Forest)")
plt.xlabel("Importance Score")
plt.tight_layout()
plt.show()

# --- Visual: One Tree from the Random Forest ---
plt.figure(figsize=(20, 10))

```

```

plot_tree(rf_clf.estimators_[0], # Plot the first tree in the forest
          feature_names=X.columns,
          class_names=['Not Attrited', 'Attrited'],
          filled=True, rounded=True, fontsize=8, )
plt.title("Example Tree from Random Forest")
plt.tight_layout()
plt.show()

```

Random Forest Classifier Results:

Accuracy: 0.9370


ROC AUC Score: 0.9747


Confusion Matrix:

```

[[1176  451]
 [ 187 8313]]

```





6.4 Boosting of Trees: AdaBoost

AdaBoost is an ensemble machine learning method that combines multiple weak learners to create a strong classifier or regressor. It works by iteratively training a sequence of models, where each new model focuses on instances incorrectly classified by previous models.

```
import pandas as pd
import numpy as np
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix, roc_curve, auc, classification_report
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import plot_tree

# --- Define X and y ---
X = df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
'Total_Revolving_Bal', 'Total_Trans_Amt']]
y = df_encoded['Attrition_Flag']

# --- Train-Test Split ---
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, stratify=y, random_state=42)

# --- Fit AdaBoost ---
ada_clf = AdaBoostClassifier(n_estimators=100, random_state=42)
ada_clf.fit(X_train, y_train)

# --- Predictions ---
y_pred = ada_clf.predict(X_test)
y_proba = ada_clf.predict_proba(X_test)[:, 1]

# --- Evaluation Metrics ---
acc = accuracy_score(y_test, y_pred)
auc_score = roc_auc_score(y_test, y_proba)
cm = confusion_matrix(y_test, y_pred)

print("=== AdaBoost Classifier Results ===")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc_score:.4f}")
print("Confusion Matrix:\n", cm)
print("\nClassification Report:\n", classification_report(y_test,
y_pred))

# --- Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Oranges', cbar=False,
xticklabels=['Pred: Existing', 'Pred: Attrited'],
yticklabels=['Actual: Existing', 'Actual: Attrited'])
```

```

plt.title('AdaBoost: Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.tight_layout()
plt.show()

# --- ROC Curve ---
fpr, tpr, _ = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, label=f'AdaBoost ROC Curve (AUC = {roc_auc:.4f})',
color='crimson')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('AdaBoost ROC Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Feature Importances ---
if hasattr(ada_clf, 'feature_importances_'):
    importances = ada_clf.feature_importances_
    feature_names = X.columns
    indices = np.argsort(importances)[::-1]

    plt.figure(figsize=(8, 5))
    sns.barplot(x=importances[indices], y=feature_names[indices],
palette='magma')
    plt.title("AdaBoost Feature Importances")
    plt.xlabel("Importance Score")
    plt.tight_layout()
    plt.show()
else:
    print("AdaBoost base estimator does not expose
feature_importances_.")

```

6.5 Boosting of Trees: XGBoost

XGBoost, short for Extreme Gradient Boosting, is a popular and powerful open-source machine learning library used for both classification and regression tasks.

Compared to AdaBoost, XGBoost is a more complex algorithm that uses gradient descent to optimize the loss function

```

import pandas as pd
import numpy as np

```

```

from xgboost import XGBClassifier, plot_tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix, roc_curve, auc, classification_report
import matplotlib.pyplot as plt
import seaborn as sns

# --- Define X and y ---
X = df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
'Total_Revolving_Bal', 'Total_Trans_Amt']]
y = df_encoded['Attrition_Flag']

# --- Train-Test Split ---
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, stratify=y, random_state=42)

# --- Fit XGBoost Classifier ---
xgb_clf = XGBClassifier(
    n_estimators=100,
    learning_rate=0.1,
    max_depth=3,
    use_label_encoder=False,
    eval_metric='logloss',
    random_state=42
)
xgb_clf.fit(X_train, y_train)

# --- Predictions ---
y_pred = xgb_clf.predict(X_test)
y_proba = xgb_clf.predict_proba(X_test)[:, 1]

# --- Evaluation ---
acc = accuracy_score(y_test, y_pred)
auc_score = roc_auc_score(y_test, y_proba)
cm = confusion_matrix(y_test, y_pred)

print("=== XGBoost Classifier Results ===")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc_score:.4f}")
print("Confusion Matrix:\n", cm)
print("\nClassification Report:\n", classification_report(y_test,
y_pred))

# --- Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Purples', cbar=False,
            xticklabels=['Pred: Existing', 'Pred: Attrited'],
            yticklabels=['Actual: Existing', 'Actual: Attrited'])
plt.title('XGBoost: Confusion Matrix')
plt.xlabel('Predicted')

```

```

plt.ylabel('Actual')
plt.tight_layout()
plt.show()

# --- ROC Curve ---
fpr, tpr, _ = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, color='darkgreen', lw=2, label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('XGBoost: ROC Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Feature Importances ---
importances = xgb_clf.feature_importances_
features = X.columns
indices = np.argsort(importances)[::-1]

plt.figure(figsize=(8, 5))
sns.barplot(x=importances[indices], y=features[indices],
palette='viridis')
plt.title("XGBoost Feature Importances")
plt.xlabel("Importance Score")
plt.tight_layout()
plt.show()

```

6.6 Support Vector Machines (3-Dimensional)

A Support Vector Machine (SVM) is a supervised learning algorithm that can be used for classification and regression tasks.

It finds the best hyperplane to separate different classes of data points, maximizing the distance between the hyperplane and the closest data points of each class.

```

import numpy as np
import pandas as pd
import plotly.graph_objects as go
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.svm import SVC
from sklearn.metrics import (
    confusion_matrix, accuracy_score, precision_score, recall_score,

```

```

f1_score
)
from sklearn.preprocessing import StandardScaler

# --- Define X and y ---
X = df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
'Total_Revolving_Bal']]
y = df_encoded['Attrition_Flag']

# --- Standardize Features ---
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# --- Train SVM with linear kernel ---
svm_clf = SVC(kernel='linear', random_state=42)
svm_clf.fit(X_scaled, y)

# --- Predict ---
y_pred = svm_clf.predict(X_scaled)

# --- Prepare DataFrame for Visualisation ---
X_3d = pd.DataFrame(X_scaled, columns=X.columns)
X_3d['Actual'] = y.values
X_3d['Predicted'] = y_pred

# --- Classification Outcome Tagging ---
def get_outcome(row):
    if row['Actual'] == 0 and row['Predicted'] == 0:
        return 'True Existing'
    elif row['Actual'] == 1 and row['Predicted'] == 1:
        return 'True Attrited'
    elif row['Actual'] == 1 and row['Predicted'] == 0:
        return 'False Existing'
    else:
        return 'False Attrited'

X_3d['Outcome'] = X_3d.apply(get_outcome, axis=1)

# --- Confusion Matrix Heatmap ---
cm = confusion_matrix(X_3d['Actual'], X_3d['Predicted'])
plt.figure(figsize=(6, 4))
sns.heatmap(
    cm, annot=True, fmt='d', cmap='YlOrRd', cbar=False,
    xticklabels=['Predicted Existing', 'Predicted Attrited'],
    yticklabels=['Actual Existing', 'Actual Attrited']
)
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted Label')
plt.ylabel('Actual Label')
plt.tight_layout()

```



```

plt.show()

# --- 3D Scatter Plot ---
color_map = {
    'True Attrited': 'green',
    'True Existing': 'red',
    'False Attrited': 'yellow',
    'False Existing': 'orange'
}

scatter = go.Scatter3d(
    x=X_3d['Total_Trans_Ct'],
    y=X_3d['Total_Ct_Chng_Q4_Q1'],
    z=X_3d['Total_Revolving_Bal'],
    mode='markers',
    marker=dict(
        size=5,
        color=[color_map[o] for o in X_3d['Outcome']],
        opacity=0.7
    ),
    text=X_3d['Outcome'],
    name='Classification Outcome'
)

# --- SVM Decision Hyperplane ---
x_range = np.linspace(X_3d['Total_Trans_Ct'].min(),
X_3d['Total_Trans_Ct'].max(), 10)
y_range = np.linspace(X_3d['Total_Ct_Chng_Q4_Q1'].min(),
X_3d['Total_Ct_Chng_Q4_Q1'].max(), 10)
xx, yy = np.meshgrid(x_range, y_range)

coef = svm_clf.coef_[0]
intercept = svm_clf.intercept_[0]
zz = (-coef[0] * xx - coef[1] * yy - intercept) / coef[2]

plane = go.Surface(
    x=xx, y=yy, z=zz,
    showscale=False,
    opacity=0.3,
    colorscale=[[0, 'red'], [1, 'red']],
    name='SVM Hyperplane'
)

# --- Combine & Render Plot ---
fig = go.Figure(data=[scatter, plane])
fig.update_layout(
    title='SVM Classification Outcome in 3D Feature
Space<br><sub>Green: True Attrited | Red: True Existing | Yellow:
False Attrited | Orange: False Existing</sub>',
    scene=dict(

```

```

        xaxis_title='Total_Trans_Ct',
        yaxis_title='Total_Ct_Chng_Q4_Q1',
        zaxis_title='Total_Revolving_Bal'
    ),
    width=1000,
    height=800,
    legend_title_text='Prediction Outcome',
    margin=dict(l=20, r=20, b=20, t=80)
)
fig.write_html("svm_3d_colored_outcome_with_heatmap.html")
fig.show()

# --- Evaluation Metrics ---
accuracy = accuracy_score(X_3d['Actual'], X_3d['Predicted'])
precision = precision_score(X_3d['Actual'], X_3d['Predicted'])
recall = recall_score(X_3d['Actual'], X_3d['Predicted'])
f1 = f1_score(X_3d['Actual'], X_3d['Predicted'])

print("SVM Performance Metrics:")
print(f"Accuracy : {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall    : {recall:.4f}")
print(f"F1 Score  : {f1:.4f}")

```

6.7 Support Vector Machine (with dimensionality reduction using PCA)

Due to the poor performance of the 3D SVM, we decided to try Dimensionality Reduction with PCA first, before performing SVM again.

Indeed, it showed much better performance than the 3D SVM.

```

import pandas as pd
import numpy as np
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix, roc_curve, auc
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns

# --- Define X and y ---
X = df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
'Total_Revolving_Bal', 'Total_Trans_Amt']]
y = df_encoded['Attrition_Flag']

# --- Standardize features ---
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

```

```

# --- Fit Support Vector Machine ---
svm_clf = SVC(kernel='rbf', probability=True, random_state=42)
svm_clf.fit(X_scaled, y)

# --- Predict ---
y_pred_class = svm_clf.predict(X_scaled)
y_pred_proba = svm_clf.predict_proba(X_scaled)[: , 1]

# --- Evaluation ---
acc = accuracy_score(y, y_pred_class)
auc_score = roc_auc_score(y, y_pred_proba)
cm = confusion_matrix(y, y_pred_class)

print("SVM Classifier Results:")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc_score:.4f}")
print("Confusion Matrix:")
print(cm)

# --- Confusion Matrix Heatmap ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Oranges', cbar=False,
            xticklabels=['Pred: Attrited', 'Pred: Existing'],
            yticklabels=['Actual: Attrited', 'Actual: Existing'])
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()

# --- ROC Curve ---
fpr, tpr, _ = roc_curve(y, y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, color='purple', lw=2, label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('SVM ROC Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

# --- PCA Reduction for 2D Plot ---
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

```

```

# Refit SVM on PCA space for decision boundary visualization
svm_pca_clf = SVC(kernel='rbf', probability=True)
svm_pca_clf.fit(X_pca, y)

# --- Decision boundary plot in 2D (PCA) ---
x_min, x_max = X_pca[:, 0].min() - 1, X_pca[:, 0].max() + 1
y_min, y_max = X_pca[:, 1].min() - 1, X_pca[:, 1].max() + 1
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 500),
                     np.linspace(y_min, y_max, 500))
Z = svm_pca_clf.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.figure(figsize=(10, 6))
plt.contourf(xx, yy, Z, alpha=0.3, cmap='coolwarm')
sns.scatterplot(x=X_pca[:, 0], y=X_pca[:, 1], hue=y, palette='Set1',
edgecolor='k')
plt.title('SVM Decision Regions (PCA-reduced 2D space)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend(title='Attrition_Flag')
plt.tight_layout()
plt.show()

```

6.8 Neural Network (MLPClassifier)

Lastly, We decided to try using a Neural Network to predict Attrition Rate.

```

import pandas as pd
import numpy as np
from sklearn.neural_network import MLPClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, roc_auc_score,
confusion_matrix, roc_curve, auc, f1_score
import matplotlib.pyplot as plt
import seaborn as sns

# Define X and y
X = df_encoded[['Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
'Total_Revolving_Bal', 'Total_Trans_Amt']]
y = df_encoded['Attrition_Flag']

# --- Step 1: Standardize features for neural networks ---
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# --- Step 2: Define and train the neural network ---
mlp = MLPClassifier(hidden_layer_sizes=(16, 8), activation='relu',
solver='adam',
                    max_iter=500, random_state=42)

```

```

mlp.fit(X_scaled, y)

# --- Step 3: Predictions ---
y_pred_class = mlp.predict(X_scaled)
y_pred_proba = mlp.predict_proba(X_scaled)[:, 1]

# --- Step 4: Evaluation metrics ---
acc = accuracy_score(y, y_pred_class)
auc_score = roc_auc_score(y, y_pred_proba)
f1 = f1_score(y, y_pred_class)
cm = confusion_matrix(y, y_pred_class)

print("Neural Network (MLPClassifier) Results:")
print(f"Accuracy: {acc:.4f}")
print(f"ROC AUC Score: {auc_score:.4f}")
print(f"F1 Score: {f1:.4f}")
print("Confusion Matrix:")
print(cm)

# --- Step 5: Heatmap of Confusion Matrix ---
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Pred: Attrited', 'Pred: Existing'],
            yticklabels=['Actual: Attrited', 'Actual: Existing'])
plt.title('Confusion Matrix Heatmap (Neural Net)')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()

# --- Step 6: ROC Curve ---
fpr, tpr, _ = roc_curve(y, y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve (Neural Network)')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

```

Conclusion

Our study delivered the following key takeaways:

- **Top churn drivers:**
 - a. **Total_Trans_Ct** (highest correlation & F-score)
 - b. **Total_Trans_Amt** (highest mutual information)
 - c. **Avg_Utilization_Ratio, Total_Revolving_Bal, Contacts_Count_12_mon.**
- **Model performance:**
 - **XGBoost** led with **94.90% accuracy** and **F1 0.8323**.
 - **MLP Classifier** close behind at **94.48% accuracy** and **F1 0.8288**.
 - SVM (with PCA) and Random Forest also delivered strong results (~ 92–93% accuracy).
- **Recommended actions:**
 - a. **Educational workshops** on account usage & budgeting.
 - b. **Credit-rebuilding programmes** for high-utilisation customers.
 - c. **Personalised retention perks** (cashback, loyalty bonuses).
 - d. **Early-warning alerts** for at-risk segments identified via clustering.

By integrating deep EDA, robust feature ranking, and ensemble modelling, this framework equips banking teams to **anticipate** and **mitigate** churn—preserving revenue and strengthening customer loyalty.